

Exact minimax risk for linear least squares, and the lower tail of sample covariance matrices

Jaouad Mourtada*

December 23, 2019

Abstract

The first part of this paper is devoted to the decision-theoretic analysis of random-design linear prediction with square loss. It is known from [Tsy03] that, under boundedness constraints on the response (and thus on regression coefficients), the minimax excess risk scales as $C\sigma^2 d/n$ up to a constant factor, where d is the model dimension, n the sample size, and σ^2 the noise parameter. Here, we study the expected excess risk with respect to the full linear class. We show that the ordinary least squares (OLS) estimator is exactly minimax optimal in the well-specified case for every distribution of covariates and noise level. Further, we express the minimax risk in terms of the distribution of *statistical leverage scores* of individual samples. We deduce a precise minimax lower bound of $\sigma^2 d/(n-d+1)$, valid for any distribution of covariates, which nearly matches the risk of OLS for centered Gaussian covariates. We then obtain nonasymptotic upper bounds on the minimax risk for covariates that satisfy a “small ball”-type regularity condition, which scale as $(1+o(1))\sigma^2 d/n$ as $d = o(n)$, both in the well-specified and misspecified cases.

Our main technical contribution is the study of the lower tail of the smallest singular value of empirical covariance matrices around 0. First, we establish a lower bound on this lower tail, valid for any distribution in dimension $d \geq 2$. We then provide a matching upper bound under a necessary regularity condition on the distribution. Our proof relies on the PAC-Bayesian technique for controlling empirical processes, and extends an analysis of [Oli16] devoted to a different part of the lower tail. Equivalently, our upper bound shows that the operator norm of the inverse sample covariance matrix has bounded L^q norm up to $q \asymp n$, and our lower bound implies that this exponent is unimprovable. Finally, we show that the regularity condition on the design naturally holds for independent coordinates.

Keywords: Least squares, covariance matrices, lower bounds, decision theory, statistical learning theory, anti-concentration.

1 Introduction

The linear least-squares problem, also called *random-design linear regression* or *linear aggregation*, is a standard problem in statistics and learning theory. Specifically, given a random pair (X, Y) where X is a covariate vector in \mathbf{R}^d and Y is a scalar response, the aim is to predict Y using a linear function $\langle \beta, X \rangle = \beta^\top X$ (with $\beta \in \mathbf{R}^d$) of X as well as possible, in a sense measured by the prediction risk with squared error $R(\beta) = \mathbb{E}[(Y - \langle \beta, X \rangle)^2]$. The best prediction is achieved by the population risk minimizer β^* , which equals:

$$\beta^* = \Sigma^{-1} \mathbb{E}[YX]$$

*MaLGa research center, Laboratory for Computational and Statistical Learning, DIBRIS, University of Genova, Italy. This work was carried while the author was a PhD student at CMAP, École polytechnique.

where $\Sigma := \mathbb{E}[XX^\top]$, assuming that both Σ and $\mathbb{E}[YX]$ are well-defined and that Σ is invertible. In the statistical setting considered here, the joint distribution P of the pair (X, Y) is unknown. The goal is then, given a sample $(X_1, Y_1), \dots, (X_n, Y_n)$ of n i.i.d. realizations of P , to find a predictor (also called *estimator*) $\hat{\beta}_n$ with small *excess risk*

$$\mathcal{E}(\hat{\beta}_n) := R(\hat{\beta}_n) - R(\beta^*) = \|\hat{\beta}_n - \beta^*\|_\Sigma^2,$$

where we define $\|\beta\|_\Sigma^2 := \langle \Sigma\beta, \beta \rangle = \|\Sigma^{1/2}\beta\|^2$. Arguably the most common procedure is the *Ordinary Least Squares* (OLS) estimator (that is, the empirical risk minimizer), defined by

$$\hat{\beta}_n^{\text{LS}} := \arg \min_{\beta \in \mathbf{R}^d} \left\{ \frac{1}{n} \sum_{i=1}^n (Y_i - \langle \beta, X_i \rangle)^2 \right\} = \hat{\Sigma}_n^{-1} \cdot \frac{1}{n} \sum_{i=1}^n Y_i X_i.$$

In this paper, we perform a precise decision-theoretic analysis of this problem, focusing on the minimax excess risk with respect to the full linear class $\mathcal{F} = \{x \mapsto \langle \beta, x \rangle : \beta \in \mathbf{R}^d\}$. This minimax perspective is relevant when little is known (or assumed) on the optimal parameter β^* . Specifically, define the *minimax excess risk* with respect to \mathcal{F} under a set \mathcal{P} of joint distributions P on (X, Y) as:

$$\inf_{\hat{\beta}_n} \sup_{P \in \mathcal{P}} \mathbb{E}[\mathcal{E}(\hat{\beta}_n)] = \inf_{\hat{\beta}_n} \sup_{P \in \mathcal{P}} \left(\mathbb{E}[R(\hat{\beta}_n)] - \inf_{\beta \in \mathbf{R}^d} R(\beta) \right), \quad (1)$$

where the infimum in (1) spans over all estimators $\hat{\beta}_n$ based on n samples, while the expectation and the risk R depend the underlying distribution P . In general, the response Y may be written as

$$Y = \langle \beta^*, X \rangle + \varepsilon, \quad (2)$$

where $\beta^* \in \mathbf{R}^d$ is the risk minimizer, which amounts to $\mathbb{E}[\varepsilon X] = \mathbb{E}[YX] - \Sigma\beta^* = 0$, and ε is the *error*. The distribution P of (X, Y) is then characterized by the distribution P_X of X , the coefficient $\beta^* \in \mathbf{R}^d$ as well as the conditional distribution of ε given X , which satisfies $\mathbb{E}[\varepsilon^2] \leq \mathbb{E}[Y^2] < +\infty$ and $\mathbb{E}[\varepsilon X] = 0$. For fixed P_X and σ^2 , we consider the following three classes \mathcal{P} of distributions, where Y is given by (2):

$$\mathcal{P}_{\text{Gauss}}(P_X, \sigma^2) := \left\{ P = P_{(X,Y)} : X \sim P_X, \beta^* \in \mathbf{R}^d, \varepsilon|X \sim \mathcal{N}(0, \sigma^2) \right\}; \quad (3)$$

$$\mathcal{P}_{\text{well}}(P_X, \sigma^2) := \left\{ P = P_{(X,Y)} : X \sim P_X, \beta^* \in \mathbf{R}^d, \mathbb{E}[\varepsilon|X] = 0, \mathbb{E}[\varepsilon^2|X] \leq \sigma^2 \right\}; \quad (4)$$

$$\mathcal{P}_{\text{mis}}(P_X, \sigma^2) := \left\{ P = P_{(X,Y)} : X \sim P_X, \beta^* \in \mathbf{R}^d, \mathbb{E}[\varepsilon^2|X] \leq \sigma^2 \right\}. \quad (5)$$

The above three classes depend on the distribution P_X of covariates, as well as the *noise level* σ^2 , which bounds the conditional second moment of the error ε . The class $\mathcal{P}_{\text{Gauss}}$ corresponds to the standard case of independent Gaussian noise, while $\mathcal{P}_{\text{well}}$ includes all *well-specified* distributions, such that the true regression function $x \mapsto \mathbb{E}[Y|X = x]$ is linear. Finally, \mathcal{P}_{mis} corresponds to the general *misspecified* case, where the regression function $x \mapsto \mathbb{E}[Y|X = x]$ is not assumed to be linear. We will study the hardness of the prediction problem (as measured by the minimax excess risk (1)) over these classes, as well as the dependence on P_X, σ^2 of the minimax risk and minimax predictor.

The second part of this paper (Section 3) is devoted to the study of the *sample covariance matrix*

$$\hat{\Sigma}_n := \frac{1}{n} \sum_{i=1}^n X_i X_i^\top, \quad (6)$$

where X_1, \dots, X_n are i.i.d. samples from P_X . In particular, we study the lower relative deviations of $\hat{\Sigma}_n$ with respect to its expectation Σ . This amounts to controlling tail probabilities of the

smallest eigenvalue $\lambda_{\min}(\tilde{\Sigma}_n)$ of the rescaled sample covariance matrix $\tilde{\Sigma}_n := \Sigma^{-1/2} \hat{\Sigma}_n \Sigma^{-1/2}$, namely to establish bounds of the form

$$\mathbb{P}(\lambda_{\min}(\tilde{\Sigma}_n) \leq t) \leq \delta, \quad (7)$$

where $t, \delta \in (0, 1)$, and their relationship depends on n, d and on P_X . The need for such bounds naturally arises in the study of the minimax excess risk, where inverse moments of $\tilde{\Sigma}_n$ (that is, quantities like $\mathbb{E}[\text{Tr}(\tilde{\Sigma}_n^{-1})]$ or $\mathbb{E}[\lambda_{\min}(\tilde{\Sigma}_n)^{-q}]$ for some $q \geq 1$) appear.

Sub-Gaussian tail bounds for $\lambda_{\min}(\tilde{\Sigma}_n)$, of the form (7) with

$$\delta = \exp\left(-cn\left(1 - C\sqrt{\frac{d}{n}} - t\right)_+^2\right)$$

for some constants c, C depending on P_X , as well as similar bounds for the largest eigenvalue $\lambda_{\max}(\tilde{\Sigma}_n)$, can be obtained under the (strong) assumption that X is sub-Gaussian (see, e.g., [Ver12]). Remarkably, it has been shown by [Oli16, KM15] that such bounds can be obtained for the *smallest* eigenvalue under much weaker assumptions on X , namely bounded fourth moments of the linear marginals of X .

While sub-Gaussian bounds establish closeness of the sample covariance matrix $\hat{\Sigma}_n$ to its empirical counterpart Σ with high (exponential) probability, they do not suffice to control negative moments of $\lambda_{\min}(\tilde{\Sigma}_n)$. Indeed, for such bounds $\delta = \delta(t)$ does not tend to 0 as $t \rightarrow 0$; in other words, they provide no nonvacuous guarantee (7) with $t = t(\delta) > 0$ as the confidence level $1 - \delta$ tends to 1. In Section 3, we complement the sub-Gaussian tail bounds by non-asymptotic large deviation bounds (7) with $\delta = \exp(-n\psi(t))$ valid uniformly over $t \in (0, c)$ for $n \geq Cd$, where the rate $\psi(t)$ tends to $+\infty$ as $t \rightarrow 0$.

Summary of results. Let us provide an overview of our results on least squares regression, which appear in Section 2:

1. We determine the exact minimax excess risk in the *well-specified case* (namely, over the classes $\mathcal{P}_{\text{well}}(P_X, \sigma^2)$ and $\mathcal{P}_{\text{Gauss}}(P_X, \sigma^2)$) for every distribution P_X of features and noise level σ^2 . For some *degenerate* distributions (Definition 1), the minimax risk is infinite (Proposition 1); while for non-degenerate ones, the OLS estimator is exactly minimax (Theorem 1) irrespective of P_X, σ^2 .
2. We express the minimax risk in terms of the distribution of *statistical leverage scores* of samples drawn from P_X (Theorem 2). Quite intuitively, distributions for which leverage scores are more uneven are seen to be harder from a minimax point of view. We deduce from this a precise minimax lower bound of $\sigma^2 d / (n - d + 1)$, valid for *every* distribution P_X of covariates. This lower bound nearly matches the $\sigma^2 d / (n - d - 1)$ risk for centered Gaussian covariates and is tight in both low and moderate dimensions; hence, Gaussian features are almost the “easiest” ones in terms of minimax risk.
3. We then turn to *upper bounds* on the minimax excess risk. Under some quantitative variant of the non-degeneracy assumption (Assumption 1) together with a fourth-moment condition on P_X (Assumption 2 or 3), we show that the minimax risk is finite and scales as $(1 + o(1))\sigma^2 d / n$ for $d = o(n)$, both in the well-specified (Theorem 3) and misspecified (Proposition 3) cases. This shows in particular that OLS is asymptotically minimax in the misspecified case as well, as $d/n \rightarrow 0$. To the best of our knowledge, these are the first bounds on the expected risk of the OLS estimator with non-Gaussian random design.

The previous upper bounds rely on the results from Section 3 on the lower tail of the sample covariance matrix $\widehat{\Sigma}_n$. Our results here are the following (assuming, to simplify notations, that $\mathbb{E}[XX^\top] = I_d$):

4. First, we establish a *lower bound* on the lower tail of $\lambda_{\min}(\widehat{\Sigma}_n)$, for $d \geq 2$ and any distribution P_X such that $\mathbb{E}[XX^\top] = I_d$, of the form: $\mathbb{P}(\lambda_{\min}(\widehat{\Sigma}_n) \leq t) \geq (ct)^{n/2}$ for some numerical constant c and every $t \in (0, 1)$ (Proposition 4). We also exhibit a “small-ball” condition (Assumption 1) which is necessary to achieve similar upper bounds.
5. Under Assumption 1, we show a matching *upper bound* on the lower tail $\mathbb{P}(\lambda_{\min}(\widehat{\Sigma}_n) \leq t)$, valid *uniformly over* $t \in (0, 1)$, and in particular for $t \rightarrow 0$. This result (Theorem 4) is the core technical contribution of this paper. Its proof relies the PAC-Bayesian technique for controlling empirical processes, which was used by [Oli16] to control a different part of the lower tail; however, some non-trivial refinements (such as non-Gaussian smoothing) are needed to handle arbitrarily small values of t . This result can be equivalently restated as an upper bound on the negative moments of $\lambda_{\min}(\widehat{\Sigma}_n)$, namely $\|\lambda_{\min}(\widehat{\Sigma}_n)^{-1}\|_{L^q} = O(1)$ for $q \asymp n$ (Corollary 4).
6. Finally, we discuss in Section 3.3 the case of independent covariates. In this case, the “small-ball” condition (Assumption 1) holds naturally under mild regularity assumptions on the distribution of individual coordinates. A result of [RV14] establishes this for coordinates with bounded density; we complement it by a general anti-concentration result for linear combination of independent variables (Proposition 6), which implies Assumption 1 under weak conditions.

Related work. Linear least squares regression is a classical statistical problem, and the literature on this topic is too vast to be surveyed here; we refer to [GKKW02, AC10, HKZ14] (and reference therein) for a more thorough overview. Analysis of least squares regression is most standard and straightforward in the *fixed design* setting, where the covariates X_1, \dots, X_n are treated as deterministic and the risk is evaluated within-sample; in this case, the expected excess risk of the OLS estimator is bounded by $\sigma^2 d/n$ (see, e.g., [Was06]).

In the random design setting considered here, a classical result ([GKKW02], Theorem 11.3) states that, if $\text{Var}(\varepsilon|X) \leq \sigma^2$ and the true regression function $g^*(x) := \mathbb{E}[Y|X = x]$ satisfies $|g^*(X)| \leq L^2$ almost surely, then the risk of the (nonlinear) *truncated* ERM estimator, defined by $\widehat{g}_n^L(x) = \min(-L, \max(L, \langle \widehat{\beta}_n^{\text{LS}}, x \rangle))$, is at most

$$\mathbb{E}[R(\widehat{g}_n^L)] - R(g^*) \leq 8(R(\beta^*) - R(g^*)) + C \max(\sigma^2, L^2) \frac{d(\log n + 1)}{n} \quad (8)$$

for some universal constant $C > 0$. This result is an *inexact oracle inequality*, where the risk is bounded by a constant times that of the best linear predictor β^* . Such guarantees are adequate in a nonparametric setting, where the approximation error $R(\beta^*) - R(g^*)$ of the linear model is itself of order $O(d/n)$ [GKKW02]. On the other hand, when no assumption is made on the magnitude of the approximation error, this bound does not ensure that the risk of the estimator approaches that of β^* . By contrast, in the *linear aggregation* problem as defined by [Nem00] (and studied by [Tsy03, Cat04, BTW07, AC11, HKZ14, LM16, Men15, Oli16]), one seeks to obtain excess risk bounds, also called *exact* oracle inequalities (where the constant 8 in the bound (8) is replaced by 1), with respect to the linear class. In this setting, Tsybakov [Tsy03] showed that the minimax rate of aggregation is of order $O(d/n)$, under boundedness assumptions on the regression function and on covariates. It is also worth noting that bounds on the regression function also implicitly constrain the optimal regression parameter to lie in some ball. This

contrasts with the approach considered here, where minimax risk with respect to the full linear class is considered. Perhaps most different from the point of view adopted here is the approach from [Fos91, Vov01, AW01, Sha15, BKM⁺15], who consider worst-case covariates (either in the individual sequences or in the agnostic learning setting) under boundedness assumptions on both covariates and outputs, and investigate achievable excess risk (or regret) bounds with respect to bounded balls in this case. By contrast, we take the distribution of covariates as given and allow the optimal regression parameter to be arbitrary, and study under which conditions on the covariates uniform bounds are achievable. Another type of non-uniform guarantees over linear classes is achieved by Ridge regression [Hoe62, Tik63] in the context of finite-dimensional or nonparametric reproducing kernel Hilbert spaces [CS02a, CS02b, DVCR05, CDV07, SZ07, SHS09, AC11, HKZ14], where the bounds do not depend explicitly on the dimension d , but rather on the spectral properties of Σ and some norm of β^* .

This work is concerned with the expected excess risk. Risk bounds in probability are obtained, among others, by [AC11, HKZ14, HS16, Oli16, Men15, LM16]. While such bounds hold with high probability, the probability is upper bounded and cannot be arbitrarily close to 1, so that they cannot be integrated to control the expected risk. Indeed, some additional regularity conditions are required in order to have finite minimax risk, as will be seen below. To the best of our knowledge, the only available uniform expected risk bounds for random-design regression are obtained in the case of Gaussian covariates, where they rely on the knowledge of the closed-form distribution of inverse covariance matrices [Ste60, BF83, And03]. One reason for considering the expected risk is that it is a single scalar, which can be more tightly controlled (in terms of matching upper and lower bounds) and compared across distributions than quantiles. In addition, random-design linear prediction is a classical statistical problem, which justifies its precise decision-theoretic analysis. On the other hand, expected risk provides little indication on the tails of the risk in the high-confidence regime: in the case of heavy-tailed noise, the OLS estimator may perform poorly, and dedicated robust estimators may be required (see, e.g., the references in [LM19]).

Another line of work [EK13, Dic16, DM16, EK18, DW18, HMRT19] considers the limiting behavior of regression procedures in the high-dimensional asymptotic regime where d, n tend to infinity at a proportional rate, with their ratio kept constant [Hub73]. The results in this setting take the form of a convergence in probability of the risk to a limit depending on the ratio d/n as well as the properties of β^* . With the notable exception of [EK18], the previous results hold under the assumption that the covariates are either Gaussian, or have a joint independence structure that leads to the same limiting behavior in high dimension. In this approach one also lets $d, n \rightarrow \infty$ while fixing some property of the parameter β^* , while here we consider non-asymptotic bounds valid for fixed n, d and uniformly over \mathbf{R}^d .

The study of spectral properties of sample covariance matrices has a rich history (see for instance [BS10, AGZ10, Tao12] and references therein); we refer to [RV10] for an overview of results (up to 2010) on the non-asymptotic control of the smallest eigenvalue of sample covariance matrices, which is the topic of Section 3. It is well-known [Ver12] that sub-Gaussian tail bounds on both the smallest and the largest eigenvalues can be obtained under sub-Gaussian assumptions on the covariates. A series of work obtained control on these quantities under weaker assumptions [ALPTJ10, MP14, SV13]. A critical observation, which has been exploited in a series of work [SV13, KM15, Oli16, Yas14, Yas15], is that the smallest eigenvalue can be controlled under much weaker tail assumptions than the largest one. Our study follows this line of work, but considers a different part of the lower tail (namely, non-vacuous deviation bounds at every probability level), which poses some additional technical difficulties. In addition, we also provide some universal lower bound on the lower tail.

Notations. Throughout this text, the transpose of an $m \times n$ real matrix A is denoted A^\top , its trace $\text{Tr}(A)$, and vectors in \mathbf{R}^d are identified with $d \times 1$ column vectors. In addition, the coordinates of a vector $x \in \mathbf{R}^d$ are indicated as superscripts: $x = (x^j)_{1 \leq j \leq d}$. We also denote $\langle x, z \rangle := x^\top z = \sum_{j=1}^d (x^j) \cdot (z^j)$ the canonical scalar product of $x, z \in \mathbf{R}^d$, and $\|x\| := \langle x, x \rangle^{1/2}$ the associated Euclidean norm. In addition, for any symmetric and positive $d \times d$ matrix A , we define the scalar product $\langle x, z \rangle_A := \langle Ax, z \rangle$ and norm $\|x\|_A := \langle Ax, x \rangle^{1/2} = \|A^{1/2}x\|$. The $d \times d$ identity matrix is denoted I_d , while $S^{d-1} = \{x \in \mathbf{R}^d : \|x\| = 1\}$ refers to the unit sphere. The smallest and largest eigenvalues of a symmetric matrix A are denoted $\lambda_{\min}(A)$ and $\lambda_{\max}(A)$ respectively; if A is positive, then $\lambda_{\max}(A) = \|A\|_{\text{op}}$ is the operator norm of A (with respect to $\|\cdot\|$), while $\lambda_{\min}(A) = \|A^{-1}\|_{\text{op}}^{-1}$.

2 Exact minimax analysis of least-squares regression

This section is devoted to the minimax analysis of the linear least-squares problem, and in particular on the dependence of its hardness on the distribution P_X of covariates. In Section 2.1, we provide the exact minimax risk and estimator in the well-specified case, namely on the class $\mathcal{P}_{\text{well}}(P_X, \sigma^2)$. In Section 2.2, we express the minimax risk in terms of the distribution of statistical leverage scores, and deduce a general lower bound. Finally, Section 2.3 provides upper bounds on the minimax risk under some regularity condition on the distribution P_X , both in the well-specified and misspecified cases.

Throughout the paper, we assume that the feature vector X satisfies $\mathbb{E}[\|X\|^2] < +\infty$, and denote $\Sigma = \mathbb{E}[XX^\top]$ its covariance matrix (by a slight but common abuse of terminology, we refer to Σ as the covariance matrix of X even when X is not centered). In addition, we assume that Σ is invertible, or equivalently that the support of X is not contained in any hyperplane; this assumption is not restrictive, since otherwise one can simply restrict to the span of the support of X (a linear subspace of \mathbf{R}^d), and merely serves to simplify notations. Then, for every distribution of Y given X such that $\mathbb{E}[Y^2] < +\infty$, the risk $R(\beta) = \mathbb{E}[(\langle \beta, X \rangle - Y)^2]$ of any $\beta \in \mathbf{R}^d$ is finite; this risk is uniquely minimized by $\beta^* = \Sigma^{-1} \mathbb{E}[YX]$, where $\mathbb{E}[YX]$ is well-defined since $\mathbb{E}[\|YX\|] \leq \mathbb{E}[Y^2]^{1/2} \mathbb{E}[\|X\|^2]^{1/2} < +\infty$ by Cauchy-Schwarz.

2.1 Minimax analysis of linear least squares

Let us start with the following definition.

Definition 1. Let $n \geq d$. The following properties of the distribution P_X are equivalent:

1. For every linear hyperplane $H \subset \mathbf{R}^d$, $\mathbb{P}(X \in H) = 0$ (equivalently, for every $\theta \in S^{d-1}$, $\mathbb{P}(\langle \theta, X \rangle = 0) = 0$);
2. The sample covariance matrix $\hat{\Sigma}_n$ is invertible almost surely;
3. The ordinary least-squares (OLS) estimator

$$\hat{\beta}_n^{\text{LS}} := \arg \min_{\beta \in \mathbf{R}^d} \sum_{i=1}^n (\langle \beta, X_i \rangle - Y_i)^2 \quad (9)$$

is uniquely defined almost surely, and equals $\hat{\beta}_n^{\text{LS}} = \hat{\Sigma}_n^{-1} n^{-1} \sum_{i=1}^n Y_i X_i$.

When either of these properties does not hold, we say that P_X is *degenerate*.

Proof. The equivalence between the second and third points is standard: the empirical risk being convex, its global minimizers are the critical points β characterized by $\widehat{\Sigma}_n \beta = n^{-1} \sum_{i=1}^n Y_i X_i$.

We now prove that the second point implies the first, by contraposition. If $\mathbb{P}(\langle \theta, X \rangle = 0) = p > 0$ for some $\theta \in S^{d-1}$, then with probability p^n , $\langle \theta, X_i \rangle = 0$ for $i = 1, \dots, n$, so that $\widehat{\Sigma}_n \theta = n^{-1} \sum_{i=1}^n \langle \theta, X_i \rangle X_i = 0$ and thus $\widehat{\Sigma}_n$ is not invertible.

Conversely, let us now show that the first point implies the second one. Note that the latter amounts to saying that X_1, \dots, X_n span \mathbf{R}^d almost surely. In particular, it suffices to show it for $n = d$, which we do by showing that, almost surely, $V_k := \text{span}(X_1, \dots, X_k)$ is of dimension k for $0 \leq k \leq d$, by induction on k . The case $k = 0$ is clear. Now, assume that $k \leq d$ and that V_{k-1} is of dimension $k-1 \leq d-1$ almost surely. Then, V_{k-1} is contained in a hyperplane of \mathbf{R}^d , and since X_k is independent of V_{k-1} , the first point implies that $\mathbb{P}(X_k \in V_{k-1}) = 0$, so that V_k is of dimension k almost surely. This concludes the proof. \square

Remark 1 (Intercept). Assume that $X = (X^j)_{1 \leq j \leq d}$, where $X^d \equiv 1$ is an intercept variable. Then, the distribution P_X is degenerate if and only if there exists $\theta = (\theta^j)_{1 \leq j < d} \in \mathbf{R}^{d-1} \setminus \{0\}$ and $c \in \mathbf{R}$ such that $\sum_{j=1}^{d-1} \theta^j X^j = c$ with positive probability. This amounts to say that (X^1, \dots, X^{d-1}) belongs to some fixed affine hyperplane of \mathbf{R}^{d-1} with positive probability.

The following result shows that non-degeneracy of the design distribution is necessary to obtain finite minimax risk.

Proposition 1 (Degenerate case). *Assume that either $n < d$, or that the distribution P_X of X is degenerate, in the sense of Definition 1. Then, the minimax excess risk with respect to the class $\mathcal{P}_{\text{Gauss}}(P_X, \sigma^2)$ is infinite.*

The fact that the minimax excess risk is infinite means that some dependence on the true parameter β^* (for instance, through its norm) is unavoidable in the expected risk of any estimator $\widehat{\beta}_n$. From now on and until the rest of this section, let us assume that the distribution P_X is non-degenerate, and that $n \geq d$. In particular, the OLS estimator is well-defined, and the empirical covariance matrix $\widehat{\Sigma}_n$ is invertible almost surely. Theorem 1 below provides the exact minimax excess risk and estimator in the well-specified case.

Theorem 1. *Assume that P_X is non-degenerate, $\sigma^2 > 0$ and $n \geq d$. The minimax risks over the classes $\mathcal{P}_{\text{well}}(P_X, \sigma^2)$ and $\mathcal{P}_{\text{Gauss}}(P_X, \sigma^2)$ coincide, and are equal to*

$$\inf_{\widehat{\beta}_n} \sup_{P \in \mathcal{P}_{\text{well}}(P_X, \sigma^2)} \mathbb{E}[\mathcal{E}_P(\widehat{\beta}_n)] = \inf_{\widehat{\beta}_n} \sup_{P \in \mathcal{P}_{\text{Gauss}}(P_X, \sigma^2)} \mathbb{E}[\mathcal{E}_P(\widehat{\beta}_n)] = \frac{\sigma^2}{n} \cdot \mathbb{E}[\text{Tr}(\widetilde{\Sigma}_n^{-1})], \quad (10)$$

where $\widetilde{\Sigma}_n = \Sigma^{-1/2} \widehat{\Sigma}_n \Sigma^{-1/2}$ is the rescaled empirical covariance matrix. In addition, the OLS estimator (9) achieves at most this risk over the class $\mathcal{P}_{\text{well}}(P_X, \sigma^2)$, and is therefore minimax optimal over the classes $\mathcal{P}_{\text{Gauss}}(P_X, \sigma^2)$ and $\mathcal{P}_{\text{well}}(P_X, \sigma^2)$ for every P_X and σ^2 .

The proof of Theorem 1 and Proposition 1 is provided in Section 4.2, and relies on simple decision-theoretic arguments. First, an upper bound (in the non-degenerate case) over $\mathcal{P}_{\text{well}}(P_X, \sigma^2)$ is obtained through the risk of the OLS estimator. Then, a matching lower bound on the minimax risk over the subclass $\mathcal{P}_{\text{Gauss}}(P_X, \sigma^2)$ is established by considering the Bayes risk under Gaussian prior on β^* , and concluding with a monotone convergence argument.

Remark 2 (Linear changes of covariates). The minimax risk is invariant under invertible linear transformations of the covariates x . This can be argued a priori, by noting that the class of linear functions of x is invariant under linear changes of variables. To see it from Theorem 1,

let $X' = AX$, where A is an invertible $d \times d$ matrix. Since $\Sigma' := \mathbb{E}[X'X'^\top]$ equals $A\Sigma A^\top$ and $\hat{\Sigma}'_n := n^{-1} \sum_{i=1}^n X'_i X'^\top_i$ equals $A\hat{\Sigma}_n A^\top$, we have

$$\hat{\Sigma}'^{-1}_n \Sigma' = ((A^\top)^{-1} \hat{\Sigma}_n^{-1} A^{-1})(A\Sigma A^\top) = (A^\top)^{-1} (\hat{\Sigma}_n^{-1} \Sigma) A^\top,$$

which is conjugate to $\hat{\Sigma}_n^{-1} \Sigma$ and hence has the same trace. By Theorem 1 (and since $\text{Tr}(\tilde{\Sigma}_n^{-1}) = \text{Tr}(\hat{\Sigma}_n^{-1} \Sigma)$), this implies that the minimax risk is the same for the covariates X and X' . In particular, the minimax risk for the design X is the same as the one for $\tilde{X} = \Sigma^{-1/2} X$.

Let us point out that the OLS estimator $\hat{\beta}_n^{\text{LS}}$ is minimax optimal for every distribution of covariates P_X and noise level σ^2 . This establishes the optimality of this procedure in a wide sense, and shows that the knowledge of neither of those properties of the distribution of (X, Y) is helpful to achieve improved risk uniformly over the linear class. On the other hand, when some additional knowledge on the optimal parameter β^* is available, OLS may no longer be optimal, and the knowledge of the noise level σ^2 may be helpful (this is for instance the case when β^* is drawn from a Gaussian prior, as in the proof of Theorem 1 in Section 4.2: the optimal estimator is then a Ridge estimator, which depends on σ^2).

Another consequence of Theorem 1 is that independent Gaussian noise is the least favorable noise structure (in terms of minimax risk) in the well-specified case for a given noise level σ^2 .

Finally, the convexity of the map $A \mapsto \text{Tr}(A^{-1})$ on positive matrices [Bha09] implies (by Jensen's inequality combined with the identity $\mathbb{E}[\tilde{\Sigma}_n] = I_d$) that the minimax risk (10) is always at least as large as $\sigma^2 d/n$, which is the minimax risk in the fixed-design case. We will however show in what follows that a strictly better lower bound can be obtained for $d \geq 2$.

2.2 Connection with leverage score and distribution-independent lower bound

In this section, we provide another expression for the minimax risk over the classes $\mathcal{P}_{\text{well}}(P_X, \sigma^2)$ and $\mathcal{P}_{\text{Gauss}}(P_X, \sigma^2)$, by relating it to the notion of *statistical leverage score* [HW78, CH88, Hub81].

Theorem 2 (Minimax risk and leverage score). *Under the assumptions of Theorem 1, the minimax risk (10) over the classes $\mathcal{P}_{\text{well}}(P_X, \sigma^2)$ and $\mathcal{P}_{\text{Gauss}}(P_X, \sigma^2)$ is equal to*

$$\inf_{\hat{\beta}_n} \sup_{P \in \mathcal{P}_{\text{Gauss}}(P_X, \sigma^2)} \mathbb{E}[\mathcal{E}_P(\hat{\beta}_n)] = \sigma^2 \cdot \mathbb{E} \left[\frac{\hat{\ell}_{n+1}}{1 - \hat{\ell}_{n+1}} \right] \quad (11)$$

where the expectation holds over an i.i.d. sample X_1, \dots, X_{n+1} drawn from P_X , and where $\hat{\ell}_{n+1}$ denotes the statistical leverage score of X_{n+1} among X_1, \dots, X_{n+1} , defined by:

$$\hat{\ell}_{n+1} = \left\langle \left(\sum_{i=1}^{n+1} X_i X_i^\top \right)^{-1} X_{n+1}, X_{n+1} \right\rangle. \quad (12)$$

The leverage score $\hat{\ell}_{n+1}$ of X_{n+1} among X_1, \dots, X_{n+1} measures the influence of the response Y_{n+1} on the associated fitted value $\hat{Y}_{n+1} = \langle \hat{\beta}_{n+1}^{\text{LS}}, X_{n+1} \rangle$: \hat{Y}_{n+1} is an affine function of Y_{n+1} , with slope $\hat{\ell}_{n+1} = \partial \hat{Y}_{n+1} / \partial Y_{n+1}$ [HW78, CH88]. Theorem 2 shows that the minimax predictive excess risk under the distribution P_X is characterized by the distribution of leverage scores of samples drawn from this distribution. Intuitively, uneven leverage scores (with some points having high leverage) imply that the estimator $\hat{\beta}_n^{\text{LS}}$ is determined by a smaller number of points, and therefore more noisy. This is consistent with the message from robust statistics that points with high leverage (typically seen as outliers) can be detrimental to the performance of the least squares estimator [HW78, CH88, Hub81], see also [RM16].

Proof of Theorem 2. By Theorem 1, the minimax risk over $\mathcal{P}_{\text{Gauss}}(P_X, \sigma^2)$ and $\mathcal{P}_{\text{well}}(P_X, \sigma^2)$ equals, letting $X_{n+1} \sim P_X$ be independent from X_1, \dots, X_n :

$$\begin{aligned}
\frac{\sigma^2}{n} \cdot \mathbb{E}[\text{Tr}(\tilde{\Sigma}_n^{-1})] &= \frac{\sigma^2}{n} \cdot \mathbb{E}[\text{Tr}(\hat{\Sigma}_n^{-1} \Sigma)] \\
&= \sigma^2 \cdot \mathbb{E}[\text{Tr}((n\hat{\Sigma}_n)^{-1} X_{n+1} X_{n+1}^\top)] \\
&= \sigma^2 \cdot \mathbb{E}[\langle (n\hat{\Sigma}_n)^{-1} X_{n+1}, X_{n+1} \rangle] \\
&= \sigma^2 \cdot \mathbb{E}\left[\frac{\langle (n\hat{\Sigma}_n + X_{n+1} X_{n+1}^\top)^{-1} X_{n+1}, X_{n+1} \rangle}{1 - \langle (n\hat{\Sigma}_n + X_{n+1} X_{n+1}^\top)^{-1} X_{n+1}, X_{n+1} \rangle}\right] \\
&= \sigma^2 \cdot \mathbb{E}\left[\frac{\hat{\ell}_{n+1}}{1 - \hat{\ell}_{n+1}}\right],
\end{aligned} \tag{13}$$

where equality (13) follows from Lemma 1 below, with $S = n\hat{\Sigma}_n$ and $v = X_{n+1}$. \square

Lemma 1. *Let S be a symmetric positive $d \times d$ matrix, and $v \in \mathbf{R}^d$. Then,*

$$\langle S^{-1}v, v \rangle = \frac{\langle (S + vv^\top)^{-1}v, v \rangle}{1 - \langle (S + vv^\top)^{-1}v, v \rangle}. \tag{14}$$

Proof. Since $S + vv^\top \succcurlyeq S$ is positive, it is invertible, and the Sherman-Morrison formula [HJ90] shows that

$$(S + vv^\top)^{-1} = S^{-1} - \frac{S^{-1}vv^\top S^{-1}}{1 + v^\top S^{-1}v},$$

so that

$$\langle (S + vv^\top)^{-1}v, v \rangle = v^\top S^{-1}v - \frac{v^\top S^{-1}vv^\top S^{-1}v}{1 + v^\top S^{-1}v} = \langle S^{-1}v, v \rangle - \frac{\langle S^{-1}v, v \rangle^2}{1 + \langle S^{-1}v, v \rangle} = \frac{\langle S^{-1}v, v \rangle}{1 + \langle S^{-1}v, v \rangle},$$

which implies that $\langle (S + vv^\top)^{-1}v, v \rangle \in [0, 1)$. Inverting this equality yields equation (14). \square

We now deduce from Theorem 2 a precise lower bound on the minimax risk (10), valid for every distribution of covariates P_X . By Proposition 1, it suffices to consider the case when $n \geq d$ and P_X is nondegenerate (since otherwise the minimax risk is infinite).

Corollary 1 (Minimax lower bound). *Under the assumptions of Theorem 1, the minimax risk (10) over $\mathcal{P}_{\text{Gauss}}(P_X, \sigma^2)$ satisfies*

$$\inf_{\hat{\beta}_n} \sup_{P \in \mathcal{P}_{\text{Gauss}}(P_X, \sigma^2)} \mathbb{E}[\mathcal{E}_P(\hat{\beta}_n)] \geq \frac{\sigma^2 d}{n - d + 1}. \tag{15}$$

Proof of Corollary 1. By Theorem 2, the minimax risk over $\mathcal{P}_{\text{Gauss}}(P_X, \sigma^2)$ writes:

$$\sigma^2 \cdot \mathbb{E}\left[\frac{\hat{\ell}_{n+1}}{1 - \hat{\ell}_{n+1}}\right] \geq \sigma^2 \cdot \frac{\mathbb{E}[\hat{\ell}_{n+1}]}{1 - \mathbb{E}[\hat{\ell}_{n+1}]}, \tag{16}$$

where the inequality follows from the convexity of the map $x \mapsto x/(1 - x) = 1 - 1/(1 - x)$ on $[0, 1)$. Now, observe that, by exchangeability of (X_1, \dots, X_{n+1}) ,

$$\begin{aligned}
\mathbb{E}[\hat{\ell}_{n+1}] &= \frac{1}{n+1} \sum_{i=1}^{n+1} \mathbb{E}\left[\left\langle \left(\sum_{i=1}^{n+1} X_i X_i^\top\right)^{-1} X_i, X_i \right\rangle\right] \\
&= \frac{1}{n+1} \mathbb{E}\left[\text{Tr}\left\{\left(\sum_{i=1}^{n+1} X_i X_i^\top\right)^{-1} \left(\sum_{i=1}^{n+1} X_i X_i^\top\right)\right\}\right] = \frac{d}{n+1}.
\end{aligned} \tag{17}$$

Plugging equation (17) into (16) yields the lower bound (15). \square

Since $n - d + 1 \geq n$, Corollary 1 implies a lower bound of $\sigma^2 d/n$. The minimax risk for linear regression has been determined under some additional boundedness assumptions on Y (and thus on β^*) by [Tsy03], showing that it scales as $\Theta(\sigma^2 d/n)$ up to numerical constants. The proof of the lower bound relies on information-theoretic arguments, and in particular on Fano's inequality [Tsy09]. Although widely applicable, such techniques often lead to loose constant factors. By contrast, the approach relying on Bayesian decision theory leading to Corollary 1 recovers the optimal leading constant, owing to the analytical tractability of the problem.

In fact, the lower bound of Corollary 1 is more precise than the $\sigma^2 d/n$ lower bound, in particular when the dimension d is commensurate to n . Indeed, in the case of centered Gaussian design, namely when $X \sim \mathcal{N}(0, \Sigma)$ for some positive matrix Σ , the risk of the OLS estimator (and thus, by Theorem 1, the minimax risk) can be computed exactly [And03, BF83], and equals

$$\mathbb{E}[\mathcal{E}_P(\hat{\beta}_n^{\text{LS}})] = \frac{\sigma^2 d}{n - d - 1}. \quad (18)$$

The distribution-independent lower bound of Corollary 1 is very close to the above whenever $n - d \gg 1$. Hence, it is almost the best possible distribution-independent lower bound on the minimax risk. This also shows that Gaussian design is almost the easiest design distribution, in terms of minimax risk. This can be understood as follows: degeneracy (a large value of $\text{Tr}(\tilde{\Sigma}_n^{-1})$) occurs whenever the rescaled sample covariance matrix $\tilde{\Sigma}_n$ is small in some direction; this occurs if either the direction of $\tilde{X} = \Sigma^{-1/2} X$ is far from uniform (so that the projection of \tilde{X} in some direction is small), or if its norm can be small. If $\tilde{X} \sim \mathcal{N}(0, I_d)$, then $\tilde{X}/\|\tilde{X}\|$ is uniformly distributed on the unit sphere, while $\|\tilde{X}\| = \sqrt{\sum_{j=1}^d (\tilde{X}^j)^2}$ is sharply concentrated around \sqrt{d} : with exponential probability, $\|\tilde{X}\| = \sqrt{d} + O(1)$ [Ver18].

In particular, in the “dense” high-dimensional regime where d and n are large and commensurate, namely $d, n \rightarrow \infty$ and $d/n \rightarrow \gamma$, the lower bound of Corollary 1 matches the minimax risk (18) in the Gaussian case, which converges to $\sigma^2 \gamma/(1 - \gamma)$. The limit $\sigma^2 \gamma/(1 - \gamma)$ is also known to be *universal* in the high-dimensional regime: for covariates with independent coordinates, the excess risk converges almost surely to this limit under mild assumptions [TV04, BS10] (note however that almost sure convergence does not imply convergence in expectation: indeed, the minimax risk may be infinite, for instance for degenerate distributions). However, the “universality” of this limit behavior is questionable [EKK11, EK18], since it relies on the independence assumption, which induces in high dimension a very specific geometry of the covariates due to the concentration of measure phenomenon [Led01, BLM13]. For instance, [EK18] obtains different limiting risks for robust regression in high dimension when considering non-independent coordinates. Setting universality aside, Corollary 1 shows that the limiting excess risk obtained in the independent case provides a *lower bound* for general design distributions.

Finally, the property of the design distribution that leads to the minimal excess risk in high dimension can be formulated simply in terms of leverage scores, using Theorem 2.

Corollary 2. *Let $(d_n)_{n \geq 1}$ be a sequence of positive integers such that $d_n/n \rightarrow \gamma \in (0, 1)$, and $(P_X^{(n)})_{n \geq 1}$ a sequence of non-degenerate distributions on \mathbf{R}^{d_n} . Assume that the minimax excess risk (10) over $\mathcal{P}_{\text{well}}(P_X^{(n)}, \sigma^2)$ converges to $\sigma^2 \gamma/(1 - \gamma)$. Then, the distribution of the leverage score $\tilde{\ell}_{n+1}^{(n)}$ of one sample among $n + 1$ under $P_X^{(n)}$ converges in probability to γ .*

Proof. Let $\phi(x) = x/(1 - x)$ for $x \in [0, 1)$, and $\psi(x) := \phi(x) - \phi(\gamma) - \phi'(\gamma)(x - \gamma)$ (with $\psi(\gamma) = 0$). Since ϕ is strictly convex, $\psi(x) > 0$ for $x \neq \gamma$, and ψ is also strictly convex. Hence, ψ is decreasing on $[0, \gamma]$ and increasing on $[\gamma, 1)$. In particular, for every $\varepsilon > 0$, $\eta_\varepsilon := \inf_{|x - \gamma| \geq \varepsilon} \psi(x) > 0$.

By Theorem 2, the assumption of Corollary 2 means that $\mathbb{E}[\phi(\tilde{\ell}_{n+1}^{(n)})] \rightarrow \phi(\gamma)$. Since in addition $\mathbb{E}[\tilde{\ell}_{n+1}^{(n)}] = d_n/(n + 1) \rightarrow \gamma$ (the first equality, used in the proof of Corollary 1, holds

for $d_n \leq n + 1$, hence for n large enough since $\gamma < 1$), we have $\mathbb{E}[\psi(\hat{\ell}_{n+1}^{(n)})] \rightarrow 0$. Now, for every $\varepsilon > 0$, $\psi(x) \geq \eta_\varepsilon \cdot \mathbf{1}(|x - \gamma| \geq \varepsilon)$, so that $\mathbb{P}(|\hat{\ell}_{n+1}^{(n)} - \gamma| \geq \varepsilon) \leq \eta_\varepsilon^{-1} \mathbb{E}[\psi(\hat{\ell}_{n+1}^{(n)})] \rightarrow 0$. \square

2.3 Upper bounds on the minimax risk

In this section, we complement the lower bound of Corollary 1 by providing matching *upper bounds* on the minimax risk. Since by Proposition 1 the minimax risk is infinite when the design distribution is degenerate, some condition is required in order to control this quantity. We therefore introduce the following quantitative version of the non-degeneracy condition:

Assumption 1 (Small-ball condition). The whitened design $\tilde{X} := \Sigma^{-1/2}X$ satisfies the following: there exist constants $C \geq 1$ and $\alpha \in (0, 1]$ such that, for every linear hyperplane H of \mathbf{R}^d and $t > 0$,

$$\mathbb{P}(\text{dist}(\tilde{X}, H) \leq t) \leq (Ct)^\alpha. \quad (19)$$

Equivalently, for every $\theta \in \mathbf{R}^d \setminus \{0\}$ and $t > 0$,

$$\mathbb{P}(|\langle \theta, X \rangle| \leq t \|\theta\|_\Sigma) \leq (Ct)^\alpha. \quad (20)$$

Note that the equivalence between (19) and (20) comes from the fact that the distance $\text{dist}(\tilde{X}, H)$ of \tilde{X} to the hyperplane H equals $|\langle \theta', \tilde{X} \rangle|$, where $\theta' \in S^{d-1}$ is a normal vector to H . Condition (20) is then recovered by letting $\theta = \Sigma^{-1/2}\theta'$ (such that $\|\theta\|_\Sigma = \|\theta'\| = 1$) and by homogeneity.

Assumption 1 states that \tilde{X} does not lie too close to any fixed hyperplane. This assumption is a strengthened variant of the “small ball” condition introduced by [KM15, Men15, LM16] in the analysis of sample covariance matrices and least squares regression, which amounts to assuming (20) for a *single* value of $t < C^{-1}$. This latter condition amounts to a uniform equivalence between the L^1 and L^2 norms of one-dimensional marginals $\langle \theta, X \rangle$ ($\theta \in \mathbf{R}^d$) of X [KM15]. Here, we require that the condition holds uniformly as $t \rightarrow 0$; the reason for this is that in order to control the minimax excess risk (10) (and thus $\mathbb{E}[\text{Tr}(\tilde{\Sigma}_n^{-1})]$), we are lead to control the lower tail of the rescaled covariance matrix $\tilde{\Sigma}_n$ at all probability levels. The study of the lower tail of $\tilde{\Sigma}_n$ (on which the results of this section rely) is deferred to Section 3. We also illustrate Assumption 1 in Section 3.3, by discussing conditions under which it holds in the case of independent coordinates.

First, Assumption 1 itself suffices to obtain an upper bound on the minimax risk of $O(\sigma^2 d/n)$, without any additional assumption on the upper tail of XX^\top (apart from integrability).

Proposition 2. *If Assumption 1 holds, then for every $P \in \mathcal{P}_{\text{well}}(P_X, \sigma^2)$, letting $C' := 3C^4 e^{1+9/\alpha}$ we have:*

$$\mathbb{E}[\mathcal{E}(\hat{\beta}_n^{\text{LS}})] \leq 2C' \cdot \frac{\sigma^2 d}{n}. \quad (21)$$

Proposition 2 (which is a consequence of Corollary 4 from Section 3.2) is optimal in terms of the rate of convergence; however, it exhibits the suboptimal $2C'$ factor in the leading term. As we show next, it is possible to obtain an optimal constant in the first-order term (as well as a second-order term of the correct order) under a modest additional assumption.

Assumption 2 (Bounded kurtosis of the norm). X satisfies $\mathbb{E}[\|\Sigma^{-1/2}X\|^4] \leq \kappa d^2$.

Remark 3. Assumption 2 amounts to $\mathbb{E}[\|\Sigma^{-1/2}X\|^4]^{1/4} \leq \kappa^{1/4} \mathbb{E}[\|\Sigma^{-1/2}X\|^2]^{1/2}$, which controls the L^4 norm of $\|\Sigma^{-1/2}X\|$ in terms of its L^2 norm. This condition is implied by the following L^2 - L^4 equivalence for one-dimensional marginals of X : for every $\theta \in \mathbf{R}^d$, $\mathbb{E}[\langle \theta, X \rangle^4]^{1/4} \leq \kappa^{1/4} \cdot \mathbb{E}[\langle \theta, X \rangle^2]^{1/2}$ (Assumption 3 below). Indeed, assuming that the latter holds, then taking $\theta =$

$\Sigma^{-1/2}e_j$ (where $(e_j)_{1 \leq j \leq d}$ denotes the canonical basis of \mathbf{R}^d), so that $\langle \theta, X \rangle$ is the j -th coordinate \tilde{X}^j of \tilde{X} , we get $\mathbb{E}[(\tilde{X}^j)^4] \leq \kappa \mathbb{E}[(\tilde{X}^j)^2]^2 = \kappa$ (since $\mathbb{E}[\tilde{X}\tilde{X}^\top] = I_d$). This implies that

$$\begin{aligned} \mathbb{E}[\|\tilde{X}\|^4] &= \mathbb{E}\left[\left(\sum_{j=1}^d (\tilde{X}^j)^2\right)^2\right] = \sum_{1 \leq j, k \leq d} \mathbb{E}[(\tilde{X}^j)^2(\tilde{X}^k)^2] \leq \sum_{1 \leq j, k \leq d} \mathbb{E}[(\tilde{X}^j)^4]^{1/2} \mathbb{E}[(\tilde{X}^k)^4]^{1/2} \\ &\leq \sum_{1 \leq j, k \leq d} \kappa^{1/2} \cdot \kappa^{1/2} = \kappa \cdot d^2, \end{aligned}$$

where the first inequality above comes from the Cauchy-Schwarz inequality. The converse is false: if \tilde{X} is uniform on $\{\sqrt{d}e_j : 1 \leq j \leq d\}$, then the first condition holds with $\kappa = 1$, while the second only holds for $\kappa \geq d$ (taking $\theta = e_1$). Hence, Assumption 2 on the upper tail of X is weaker than an L^2 - L^4 equivalence of the one-dimensional marginals of X ; on the other hand, we do require a small-ball condition (Assumption 1) on the lower tail of X .

Theorem 3 (Upper bound in the well-specified case). *Grant Assumptions 1 and 2. Let $C' = 3C^4e^{1+9/\alpha}$ (which only depends on α, C). If $n \geq \min(6\alpha^{-1}d, 12\alpha^{-1}\log(12\alpha^{-1}))$, then*

$$\frac{1}{n} \mathbb{E}[\text{Tr}(\tilde{\Sigma}_n^{-1})] \leq \frac{d}{n} + 8C'\kappa \left(\frac{d}{n}\right)^2. \quad (22)$$

In particular, the minimax excess risk over the class $\mathcal{P}_{\text{well}}(P_X, \sigma^2)$ satisfies:

$$\frac{\sigma^2 d}{n} \leq \inf_{\hat{\beta}_n} \sup_{P \in \mathcal{P}_{\text{well}}(P_X, \sigma^2)} \mathbb{E}[\mathcal{E}_P(\hat{\beta}_n)] \leq \frac{\sigma^2 d}{n} \left(1 + 8C' \frac{\kappa d}{n}\right). \quad (23)$$

The proof of Theorem 3 is given in Section 4.3; it relies in particular on Lemma 7 herein and on Theorem 4 from Section 3. As shown by the lower bound (established in Corollary 1), the constant in the first-order term in (23) is tight; in addition, one could see from a higher-order expansion (under additional moment assumptions) that the second-order term is also tight, up to the constant $8C'$ factor. This suggests that Assumption 2 is essentially a minimal condition on the upper tail of XX^\top to obtain a second-order term in $O((d/n)^2)$.

Let us now consider the general misspecified case, namely the class $\mathcal{P}_{\text{mis}}(P_X, \sigma^2)$. Here, we will need the slightly stronger Assumption 3.

Assumption 3 (L^2 - L^4 norm equivalence). There exists a constant $\kappa > 0$ such that, for every $\theta \in \mathbf{R}^d$, $\mathbb{E}[\langle \theta, X \rangle^4] \leq \kappa \cdot \mathbb{E}[\langle \theta, X \rangle^2]^2$.

Proposition 3 (Upper bound in the misspecified case). *Assume that P_X satisfies Assumptions 1 and 3, and that*

$$\chi := \mathbb{E}[\mathbb{E}[\varepsilon^2|X]^2 \|\Sigma^{-1/2}X\|^4] / d^2 < +\infty$$

(note that $\chi \leq \mathbb{E}[(Y - \langle \beta^*, X \rangle)^4 \|\Sigma^{-1/2}X\|^4] / d^2$). Then, for $n \geq \max(96, 6d)/\alpha$, the risk of the OLS estimator satisfies

$$\mathbb{E}[\mathcal{E}(\hat{\beta}_n^{\text{LS}})] \leq \frac{1}{n} \mathbb{E}[(Y - \langle \beta^*, X \rangle)^2 \|\Sigma^{-1/2}X\|^2] + 276C'^2 \sqrt{\kappa\chi} \left(\frac{d}{n}\right)^{3/2}. \quad (24)$$

In particular, we have

$$\frac{\sigma^2 d}{n} \leq \inf_{\hat{\beta}_n} \sup_{P \in \mathcal{P}_{\text{mis}}(P_X, \sigma^2)} \mathbb{E}[\mathcal{E}(\hat{\beta}_n)] \leq \sup_{P \in \mathcal{P}_{\text{mis}}(P_X, \sigma^2)} \mathbb{E}[\mathcal{E}(\hat{\beta}_n^{\text{LS}})] \leq \frac{\sigma^2 d}{n} \left(1 + 276C'^2 \kappa \sqrt{\frac{d}{n}}\right). \quad (25)$$

The proof of Proposition 3 is provided in Section 4.4; it combines the results from Section 3 with a tail bound from [Oli16]. Proposition 3 shows that, under Assumptions 1 and 3, the minimax excess risk over the class $\mathcal{P}_{\text{mis}}(P_X, \sigma^2)$ scales as $(1 + o(1))\sigma^2 d/n$ as $d/n \rightarrow 0$. It also shows that the OLS estimator is asymptotically minimax on the misspecified class $\mathcal{P}_{\text{mis}}(P_X, \sigma^2)$ as $d = o(n)$, and that independent Gaussian noise is asymptotically the least favorable structure for the error ε .

2.4 Parameter estimation

Let us briefly discuss how the results of this section obtained for prediction can be adapted to the problem of parameter estimation, where the loss of an estimate $\hat{\beta}_n$ given β^* is $\|\hat{\beta}_n - \beta^*\|^2$.

By the same proof as that of Theorem 1 (but replacing the norm $\|\cdot\|_\Sigma$ by $\|\cdot\|$), the minimax excess risk over the classes $\mathcal{P}_{\text{Gauss}}(P_X, \sigma^2)$ and $\mathcal{P}_{\text{well}}(P_X, \sigma^2)$ is $(\sigma^2/n)\mathbb{E}[\text{Tr}(\hat{\Sigma}_n^{-1})]$, achieved by the OLS estimator. By convexity of $A \mapsto \text{Tr}(A^{-1})$ over positive matrices [Bha09], this quantity is lower bounded by $\sigma^2 \text{Tr}(\Sigma^{-1})/n$.

In the case of centered Gaussian covariates, $\mathbb{E}[\text{Tr}(\hat{\Sigma}_n^{-1})] = \text{Tr}(\Sigma^{-1}\mathbb{E}[\tilde{\Sigma}_n^{-1}]) = \text{Tr}(\Sigma^{-1})n/(n-d-1)$ [And03], so that the minimax risk is $\sigma^2 \text{Tr}(\Sigma^{-1})/(n-d-1)$. On the other hand, an improved lower bound for general design, like that of Corollary 1 for prediction, does not appear to hold for estimation. The reason for this is that the map $A \mapsto A/(1 - \text{Tr}(A))$ is not convex over positive matrices for $d \geq 2$ (where convexity is defined with respect to the positive definite order, see e.g. [BV04] for a definition), although its trace is.

Finally, the results of Section 3 on the lower tail of $\tilde{\Sigma}_n$ can be used to obtain upper bounds in a similar fashion as for prediction. For instance, an analogue of Proposition 2 can be directly obtained by bounding $\text{Tr}(\hat{\Sigma}_n^{-1}) \leq \lambda_{\min}(\tilde{\Sigma}_n)^{-1} \cdot \text{Tr}(\Sigma^{-1})$. Since the main focus of this paper is on prediction, we do not elaborate further in this direction.

3 Bounding the lower tail of a sample covariance matrix at all probability levels

Throughout this section, up to replacing X by $\Sigma^{-1/2}X$, we assume unless otherwise stated that $\mathbb{E}[XX^\top] = I_d$. Our aim is to obtain non-asymptotic large deviation inequalities of the form:

$$\mathbb{P}(\lambda_{\min}(\hat{\Sigma}_n) \leq t) \leq e^{-n\psi(t)}$$

where $\psi(t) \rightarrow \infty$ as $t \rightarrow 0^+$. Existing bounds [Ver12, SV13, KM15, Oli16] are typically sub-Gaussian bounds with $\psi(t) = c(1 - C\sqrt{d/n} - t)_+^2$ for some constants $c, C > 0$, which become suboptimal as $t \rightarrow 0$. In this section, we study the behavior of the large deviations for small values of t (say, $t \in (0, c)$, where $c < 1$ is a fixed constant). In Section 3.1, we provide a lower bound on these tail probabilities, namely an upper bound on ψ , valid for every distribution of X when $d \geq 2$. In Section 3.2, we show that Assumption 1 is necessary and sufficient to obtain tail bounds of the optimal order. Finally, in Section 3.3 we show that Assumption 1 is naturally satisfied in the case of independent coordinates, under a mild regularity condition on their distributions.

3.1 A general lower bound on the lower tail

First, Proposition 4 below shows that in dimension $d \geq 2$, the probability of deviations of $\lambda_{\min}(\hat{\Sigma}_n)$ cannot be arbitrarily small.

Proposition 4. *Assume that $d \geq 2$. Let X be a random vector in \mathbf{R}^d such that $\mathbb{E}[XX^\top] = I_d$. Then, for every $t \leq 1$,*

$$\sup_{\theta \in S^{d-1}} \mathbb{P}(|\langle \theta, X \rangle| \leq t) \geq 0.16 \cdot t, \quad (26)$$

and therefore

$$\mathbb{P}(\lambda_{\min}(\hat{\Sigma}_n) \leq t) \geq (0.025 \cdot t)^{n/2}. \quad (27)$$

The assumption $d \geq 2$ is necessary since for $d = 1$, if $X = 1$ almost surely, then $\lambda_{\min}(\hat{\Sigma}_n) = 1$ almost surely. Proposition 4 is proved by a probabilistic argument, namely by considering a random vector θ drawn uniformly on the sphere S^{d-1} .

Proof of Proposition 4. Let Θ be a random variable distributed uniformly on the unit sphere S^{d-1} and independent of X . We have

$$\sup_{\theta \in S^{d-1}} \mathbb{P}(|\langle \theta, X \rangle| \leq t) \geq \mathbb{E}[\mathbb{P}(|\langle \Theta, X \rangle| \leq t | \Theta)] = \mathbb{E}[\mathbb{P}(|\langle \Theta, X \rangle| \leq t | X)].$$

Next, note that for every $x \in \mathbf{R}^d$, $\langle \Theta, x \rangle$ is distributed as $\|x\| \cdot \Theta_1$, where Θ_1 denotes the first coordinate of Θ . Since X is independent of Θ , the above inequality becomes

$$\sup_{\theta \in S^{d-1}} \mathbb{P}(|\langle \theta, X \rangle| \leq t) \geq \mathbb{E}\left[\mathbb{P}\left(|\Theta_1| \leq \frac{t}{\|X\|} \middle| X\right)\right]. \quad (28)$$

Now, since $\mathbb{E}[\|X\|^2] = \text{Tr}(\mathbb{E}[XX^\top]) = d$, Markov's inequality implies that $\mathbb{P}(\|X\| \geq 2\sqrt{d}) \leq \mathbb{E}[\|X\|^2]/(4d) \leq 1/4$. Since $r \mapsto \mathbb{P}_\theta(|\theta_1| \leq t/r)$ is non-increasing, plugging this into (28) yields

$$\sup_{\theta \in S^{d-1}} \mathbb{P}(|\langle \theta, X \rangle| \leq t) \geq \frac{3}{4} \cdot \mathbb{P}\left(|\Theta_1| \leq \frac{t}{2\sqrt{d}}\right). \quad (29)$$

Let us now derive the distribution of $|\Theta_1|$. Let $\phi : S^{d-1} \rightarrow \mathbf{R}$ be the projection on the first coordinate: $\phi(\theta) = \theta_1$ for $\theta \in S^{d-1}$. Note that for $u \in [-1, 1]$, $\phi^{-1}(u) = \{u\} \times (\sqrt{1-u^2} \cdot S^{d-2})$ which is isometric to $\sqrt{1-u^2} \cdot S^{d-2}$ and hence has $(d-2)$ -dimensional Hausdorff measure $C_d(1-u^2)^{(d-2)/2}$ for some constant C_d . In addition, since $\phi(\theta) = \langle e_1, \theta \rangle$ (where $e_1 = (1, 0, \dots, 0)$), $\nabla \phi(\theta) \in (\mathbf{R}\theta)^\perp$ is the orthogonal projection of e_1 on $(\mathbf{R}\theta)^\perp$, namely $e_1 - \theta_1\theta$, with norm $\|\nabla \phi(\theta)\| = \sqrt{1-\theta_1^2}$. Fix $t \in (0, 1]$ and define $g(\theta) = \mathbf{1}(|\theta_1| \leq t)/\sqrt{1-\theta_1^2}$, which equals $\mathbf{1}(|u| \leq t)/\sqrt{1-u^2}$ on $\phi^{-1}(u)$ (for $u \in (-1, 1)$), and such that $g(\theta) \cdot \|\nabla \phi(\theta)\| = \mathbf{1}(|\theta_1| \leq t)$. Hence, the coarea formula [Fed96, Theorem 3.2.2] implies that, for every $t \in (0, 1]$,

$$\begin{aligned} \mathbb{P}(|\Theta_1| \leq t) &= \int_{S^{d-1}} g(\theta) \|\nabla \phi(\theta)\| \pi(d\theta) = \int_{-1}^1 \frac{\mathbf{1}(|u| \leq t)}{\sqrt{1-u^2}} \times C_d(1-u^2)^{(d-2)/2} du \\ &= 2C_d \int_0^t (1-u^2)^{(d-3)/2} du. \end{aligned} \quad (30)$$

If $d = 2$, (30) implies that $|\Theta_1|$ has density $(2/\pi)/\sqrt{1-t^2} \geq 2/\pi$ on $[0, 1]$, and hence for $t \in [0, 1]$:

$$\mathbb{P}\left(|\Theta_1| \leq \frac{t}{2\sqrt{d}}\right) \geq \frac{2}{\pi} \times \frac{t}{2\sqrt{2}}. \quad (31)$$

If $d = 3$, (30) implies that $|\Theta_1|$ is uniformly distributed on $[0, 1]$, so that for $t \in [0, 1]$

$$\mathbb{P}\left(|\Theta_1| \leq \frac{t}{2\sqrt{d}}\right) = \frac{t}{2\sqrt{3}}. \quad (32)$$

Now, assume that $d \geq 4$. Letting $t = 1$ in (30) yields the value of the constant C_d , which normalizes the right-hand side: since $1 - u^2 \leq e^{-u^2}$,

$$\begin{aligned} (2C_d)^{-1} &= \int_0^1 (1-u^2)^{(d-3)/2} du \leq \int_0^1 e^{-(d-3)u^2/2} du \\ &\leq \frac{1}{\sqrt{d-3}} \int_0^{\sqrt{d-3}} e^{-u^2/2} du \leq \frac{1}{\sqrt{d-3}} \times \sqrt{\frac{\pi}{2}}, \end{aligned}$$

so that $2C_d \geq \sqrt{2(d-3)/\pi}$. Finally, if $u \leq 1/(2\sqrt{d})$, then

$$(1-u^2)^{(d-3)/2} \geq \left(1 - \frac{1}{4d}\right)^{d/2} \geq \left(1 - \frac{1}{16}\right)^2,$$

using the fact that $4d \geq 16$ and that the function $x \mapsto (1 - 1/x)^{x/8}$ is increasing on $(1, +\infty)$. Plugging the above lower bounds in (30) shows that, for $t \leq 1$,

$$\mathbb{P}\left(|\Theta_1| \leq \frac{t}{2\sqrt{d}}\right) = 2C_d \int_0^{t/(2\sqrt{d})} (1 - u^2)^{(d-3)/2} du \geq \sqrt{\frac{2(d-3)}{\pi}} \times \left(\frac{15}{16}\right)^2 \frac{t}{2\sqrt{d}} \geq \frac{t}{3} \quad (33)$$

where the last inequality is obtained by noting that $(d-3)/d \geq 1/4$ for $d \geq 4$ and lower bounding the resulting constant. The bounds (31), (32) and (33) imply that, for every $d \geq 2$ and $t \leq 1$,

$$\mathbb{P}\left(|\Theta_1| \leq \frac{t}{2\sqrt{d}}\right) \geq \frac{t}{\pi\sqrt{2}}. \quad (34)$$

The first inequality of Proposition 4 follows by combining inequalities (29) and (34). The second inequality (27) is a consequence of the first, by Lemma 2 below. \square

Lemma 2. *For $t \in (0, 1)$, let $p_t = \sup_{\theta \in S^{d-1}} \mathbb{P}(|\langle \theta, X \rangle| \leq t)$. Then, $\mathbb{P}(\lambda_{\min}(\hat{\Sigma}_n) \leq t) \geq p_{\sqrt{t}}^n$.*

Proof of Lemma 2. Let $p < p_{\sqrt{t}}$. By definition of $p_{\sqrt{t}}$, there exists $\theta \in S^{d-1}$ such that $\mathbb{P}(\langle \theta, X \rangle^2 \leq t) \geq p$. Hence, by independence, with probability at least p^n , $\langle \theta, X_i \rangle^2 \leq t$ for $i = 1, \dots, n$, so that $\lambda_{\min}(\hat{\Sigma}_n) \leq \langle \hat{\Sigma}_n \theta, \theta \rangle \leq t$. Taking $p \rightarrow p_{\sqrt{t}}$ concludes the proof. \square

Proposition 4 shows that $\mathbb{P}(\lambda_{\min}(\hat{\Sigma}_n) \leq t)$ is at least $(Ct)^{cn}$, where $C = 0.025$ and $c = 1/2$ are absolute constants; this bound writes $e^{-n\psi(t)}$, where $\psi(t) \asymp \log(1/t)$ as $t \rightarrow 0^+$. In the following section, we address the question of obtaining matching upper bounds on this lower tail.

3.2 Optimal control of the lower tail

In this section, we study conditions under which an upper bound matching the lower bound from Proposition 4 can be obtained. We start by noting that Assumption 1 is necessary to obtain such bounds:

Remark 4 (Necessity of the small ball condition). Assume that there exists $c_1, c_2 > 0$ such that $\mathbb{P}(\lambda_{\min}(\hat{\Sigma}_n) \leq t) \leq (c_1 t)^{c_2 n}$ for all $t \in (0, 1)$. Then, Lemma 2 implies that $p_t \leq (c_1 t^2)^{c_2}$ for all $t \in (0, 1)$. This amounts to saying that P_X satisfies Assumption 1 with $C = \sqrt{c_1}$ and $\alpha = 2c_2$.

As Theorem 4 shows, Assumption 1 is also sufficient to obtain an optimal control on the lower tail.

Theorem 4. *Let X be a random vector in \mathbf{R}^d . Assume that $\mathbb{E}[XX^\top] = I_d$, and that X satisfies Assumption 1. Then, $n \geq 6d/\alpha$, we have for every $t \in (0, 1)$:*

$$\mathbb{P}(\lambda_{\min}(\hat{\Sigma}_n) \leq t) \leq (C't)^{\alpha n/6} \quad (35)$$

where $C' = 3C^4 e^{1+9/\alpha}$.

Note that Theorem 4 can be stated in the non-isotropic case, where $\Sigma = \mathbb{E}[XX^\top]$ is arbitrary:

Corollary 3. *Let X be a random vector in \mathbf{R}^d such that $\mathbb{E}[\|X\|^2] < +\infty$, and let $\Sigma = \mathbb{E}[XX^\top]$. Assume that X satisfies Assumption 1. Then, if $d/n \leq \alpha/6$, for every $t \in (0, 1)$, the empirical covariance matrix $\hat{\Sigma}_n$ formed with an i.i.d. sample of size n satisfies*

$$\hat{\Sigma}_n \succcurlyeq t\Sigma \quad (36)$$

with probability at least $1 - (C't)^{\alpha n/6}$, where C' is as in Theorem 4.

Proof of Corollary 3. We may assume that Σ is invertible: otherwise, we can just consider the span of the support of X , which is a subspace of \mathbf{R}^d of dimension $d' \leq d \leq \alpha n/6$. Now, let $\tilde{X} = \Sigma^{-1/2}X$; by definition, $\mathbb{E}[\tilde{X}\tilde{X}^\top] = I_d$, and \tilde{X} satisfies Assumption 1 since X does. By Theorem 4, with probability at least $1 - (C't)^{\alpha n/6}$, $\lambda_{\min}(\Sigma^{-1/2}\hat{\Sigma}_n\Sigma^{-1/2}) \geq t$, which amounts to $\Sigma^{-1/2}\hat{\Sigma}_n\Sigma^{-1/2} \succcurlyeq tI_d$, and thus $\hat{\Sigma}_n \succcurlyeq t\Sigma$. \square

It is worth noting that Theorem 4 does not require any condition on the upper tail of XX^\top , aside from the assumption $\mathbb{E}[XX^\top] = I_d$. Indeed, as noted in Remark 4, it only requires the necessary Assumption 1. In particular, it does not require any sub-Gaussian assumption on X , similarly to the results from [KM15, Oli16, Yas14, Yas15]; this owes to the fact that a sum of independent positive random variables is naturally bounded away from 0.

Remark 5 (Extension to random quadratic forms). Theorem 4 directly extends (up to slight changes in the notations and proofs) to random quadratic forms $v \mapsto \langle A_i v, v \rangle$ where A_1, \dots, A_n are positive semi-definite and i.i.d., with $\mathbb{E}[A_i] = I_d$ (Theorem 4 corresponds to the rank 1 case where $A_i = X_i X_i^\top$). On the other hand, the lower bound of Proposition 4 is specific to rank 1 matrices, as can be seen by considering the counterexample where $A_i = I_d$ almost surely.

Idea of the proof. The proof of Theorem 4 is provided in Section 5. It builds on the analysis of [Oli16], who obtains sub-Gaussian deviation bounds under fourth moment assumptions (Assumption 3), although some refinements are needed to handle our considered regime (with t arbitrarily small).

The proof starts with the representation of $\lambda_{\min}(\hat{\Sigma}_n)$ as the infimum of an empirical process:

$$\lambda_{\min}(\hat{\Sigma}_n) = \inf_{\theta \in S^{d-1}} \langle \hat{\Sigma}_n \theta, \theta \rangle = \inf_{\theta \in S^{d-1}} \left\{ Z(\theta) := \frac{1}{n} \sum_{i=1}^n \langle \theta, X_i \rangle^2 \right\}. \quad (37)$$

In order to control this infimum, a natural approach is to first control $Z(\theta)$ on a suitable finite ε -covering of S^{d-1} using Assumption 1, independence, and a union bound, and then to extend this control to S^{d-1} by approximation. However, this approach (see e.g. [Ver12] for a use of this argument) fails here, since the control of the approximation term would require an exponential upper bound on $\|\hat{\Sigma}_n\|_{\text{op}}$, which requires a sub-Gaussian assumption on X . Instead, as in [Oli16], we use the so-called PAC-Bayesian technique [McA99b, McA99a, LST03, Cat07, AC11]. This technique enables one to control a smoothed version of the process $Z(\theta)$, namely

$$Z(\rho) := \int_{\mathbf{R}^d} \left(\frac{1}{n} \sum_{i=1}^n \langle \theta, X_i \rangle^2 \right) \rho(d\theta),$$

uniformly over all smoothing distributions ρ on \mathbf{R}^d whose relative entropy $\text{KL}(\rho, \pi)$ with respect to a fixed “prior” distribution π on \mathbf{R}^d is bounded. The proof then involves controlling (i) the Laplace transform of the process; (ii) the approximation term; and (iii) the entropy term. In order to control the last two, a careful choice of the smoothing distribution (and prior) is needed.

Remark 6 (PAC-Bayes vs. ε -net argument). As indicated before, the use of an ε -net argument would fail here, since it would lead to an approximation term depending on $\|\hat{\Sigma}_n\|_{\text{op}}$. On the other hand, the use of a smoothing distribution which is “isotropic” and centered at a point θ enables one to obtain an approximation term in terms of $\text{Tr}(\hat{\Sigma}_n)/d$, which can be bounded after proper truncation of X (in a way that does not overly degrade Assumption 1).

Remark 7 (Choice of prior and posteriors: entropy term). The PAC-Bayesian technique is classically employed in conjunction with Gaussian prior and smoothing distribution [LST03, AC11,

Oli16]. This choice is convenient, since both the approximation and entropy term have closed-form expressions (in addition, a Gaussian distribution centered at θ yields the desired “isotropic” approximation term).

However, in order to obtain non-vacuous bounds for arbitrarily small t , we need the approximation term (and thus the “diameter” γ of the smoothing distribution) to be arbitrarily small. But as $\gamma \rightarrow 0$, the entropy term for Gaussian distributions grows too rapidly (as d/γ^2 , instead of the $d \log(1/\gamma)$ rate suggested by covering numbers), which ultimately leads to vacuous bounds. In order to bypass this difficulty, we employ a more refined choice of prior and smoothing distributions, which leads to an optimal entropy term of $d \log(1/\gamma)$. In addition, as shown by symmetry considerations, this choice of smoothing also leads to an “isotropic” approximation term controlled by $\text{Tr}(\hat{\Sigma}_n)/d$ instead of $\|\hat{\Sigma}_n\|_{\text{op}}$.

Formulation in terms of moments. The statements of this section on the lower tail of $\lambda_{\min}(\hat{\Sigma}_n)$ can equivalently be rephrased in terms of its negative moments. For $q \geq 1$, we denote $\|Z\|_{L^q} := \mathbb{E}[|Z|^q]^{1/q} \in [0, +\infty]$ the L^q norm of a real random variable Z .

Corollary 4. *Under the assumptions of Theorem 4 and for $n \geq 12/\alpha$, then for any $1 \leq q \leq \alpha n/12$,*

$$\|\max(1, \lambda_{\min}(\hat{\Sigma}_n)^{-1})\|_{L^q} \leq 2^{1/q} \cdot C'. \quad (38)$$

Conversely, the previous inequality implies that $\mathbb{P}(\lambda_{\min}(\hat{\Sigma}_n) \leq t) \leq (2C')^{\alpha n/12}$ for all $t \in (0, 1)$.

Finally, for any random vector X in \mathbf{R}^d , $d \geq 2$, such that $\mathbb{E}[XX^\top] = I_d$, we have for any $q \geq n/2$:

$$\|\lambda_{\min}(\hat{\Sigma}_n)^{-1}\|_{L^q} = +\infty.$$

Corollary 4 directly follows from Theorem 4, Proposition 4 and Lemma 3 below.

Lemma 3. *Let Z be a nonnegative real variable.*

1. *If there exist some constants $C \geq 1$ and $a \geq 2$ such that $\mathbb{P}(Z \leq t) \leq (Ct)^a$ for all $t > 0$, then $\|Z^{-1}\|_{L^q} \leq \|\max(1, Z^{-1})\|_{L^q} \leq 2^{1/q}C \leq 2C$ for all $1 \leq q \leq a/2$.*
2. *Conversely, if $\|Z^{-1}\|_{L^q} \leq C$ for some constants $q \geq 1$ and $C > 0$, then $\mathbb{P}(Z \leq t) \leq (Ct)^q$ for all $t > 0$.*
3. *Finally, if there exist constants $c, a > 0$ such that $\mathbb{P}(Z \leq t) \geq (ct)^a$ for all $t \in (0, 1)$, then $\|Z^{-1}\|_{L^q} = +\infty$ for $q \geq a$.*

Proof. For the first point, since $\max(1, Z^{-q})$ is nonnegative, we have

$$\mathbb{E}[\max(1, Z^{-q})] = \int_0^\infty \mathbb{P}(\max(1, Z^{-q}) \geq u) du = \int_0^\infty \mathbb{P}(\min(1, Z) \leq u^{-1/q}) du.$$

For $u \leq C^q$, we bound $\mathbb{P}(\min(1, Z) \leq u^{-1/q}) \leq 1$, while for $u \geq C^q$ (so that $u^{-1/q} \leq C^{-1} \leq 1$), we bound $\mathbb{P}(\min(1, Z) \leq u^{-1/q}) = \mathbb{P}(Z \leq u^{-1/q}) \leq (Cu^{-1/q})^a$. We then conclude that

$$\|\max(1, Z^{-1})\|_{L^q}^q \leq C^q + \int_{C^q}^\infty (C^{-q}u)^{-a/q} du = C^q \left[1 + \int_1^\infty v^{-a/q} dv \right] \leq 2C^q,$$

where we let $v = C^{-q}u$ and used the fact that $\int_1^\infty v^{-a/q} dv \leq \int_1^\infty v^{-2} dv = 1$ since $q \leq a/2$. The second point follows from Markov’s inequality: for every $t > 0$,

$$\mathbb{P}(Z \leq t) = \mathbb{P}(Z^{-q} \geq t^{-q}) \leq t^q \cdot \mathbb{E}[Z^{-q}] \leq (Ct)^q.$$

Finally, for the third point, since $\mathbb{P}(Z \leq u^{-1/q}) \geq (cu^{-1/q})^a$ for $u > 1$, we have for $q \geq a$:

$$\mathbb{E}[Z^{-q}] = \int_0^\infty \mathbb{P}(Z \leq u^{-1/q}) du \geq \int_1^\infty c^a u^{-a/q} du \geq c^a \int_1^\infty u^{-1} du = +\infty. \quad \square$$

3.3 The small-ball condition for independent covariates

We now discuss conditions under which the small-ball condition (Assumption 1) holds in the case of independent coordinates. In this section, we assume that the coordinates X^j , $1 \leq j \leq d$, of $X = \tilde{X}$ are independent. Note that the condition $\mathbb{E}[XX^\top] = I_d$ means that the X^j are centered and with unit variance.

Let us introduce the *Lévy concentration function* $Q_Z : \mathbf{R}^+ \rightarrow [0, 1]$ of a real random variable Z defined by, for $t \geq 0$,

$$Q_Z(t) := \sup_{a \in \mathbf{R}} \mathbb{P}(|Z - a| \leq t).$$

Anti-concentration (or small ball probabilities) [NV13] refers to nonvacuous upper bounds on this function. Here, in order to establish Assumption 1, it suffices to show that $Q_{\langle \theta, X \rangle}(t) \leq (Ct)^\alpha$ for all $t > 0$ and $\theta \in S^{d-1}$. We are thus lead to establish anti-concentration of linear combinations of independent variables $\langle \theta, X \rangle = \sum_{j=1}^d \theta^j X^j$, uniformly over $\theta \in S^{d-1}$, namely to provide upper bounds on:

$$Q_X(t) := \sup_{\theta \in S^{d-1}} Q_{\langle \theta, X \rangle}(t).$$

Small-ball probabilities naturally appear in the study of the smallest singular value of a random matrix (see [RV10]). [TV09a, TV09b, RV08, RV09] studied anti-concentration for variables of the form $\langle \theta, X \rangle$, and deduced estimates of the smallest singular value of random matrices. These bounds are however slightly different from the one we need: indeed, they hold for “unstructured” vectors θ (which do not have additive structure, see [RV10]), rather than uniformly over $\theta \in S^{d-1}$. Here, in order to show that Assumption 1 holds, we need bounds over Q_X , which requires some additional assumption on the distribution of the coordinates X^j .

Clearly, $Q_X \geq \max_{1 \leq j \leq d} Q_{X^j}$, and in particular we need the coordinates X^j themselves to exhibit anti-concentration. Remarkably, a result of [RV14] (building on a reduction by [Rog87] to uniform variables) shows that, if the X^j have bounded densities, a reverse inequality holds:

Proposition 5 ([RV14], Theorem 1.2). *Assume that X^1, \dots, X^d are independent and have density bounded by $C_0 > 0$. Then, for every $\theta \in \mathbf{R}^d$, $\sum_{j=1}^d \theta^j X^j$ has density bounded by $\sqrt{2} C_0$. In other words, $Q_X(t) \leq 2\sqrt{2} C_0 t$ for every $t > 0$, i.e., Assumption 1 holds with $\alpha = 1$ and $C = 2\sqrt{2} C_0$.*

Equivalently, if $\max_{1 \leq j \leq d} Q_{X^j}(t) \leq Ct$ for all $t > 0$, then $Q_X(t) \leq \sqrt{2} Ct$ for all $t > 0$, and the constant $\sqrt{2}$ is optimal [RV14]. Whether a general bound of Q_X in terms of $\max_{1 \leq j \leq d} Q_{X^j}$ holds is unclear (for instance, the inequality $Q_X \leq \sqrt{2} \max_{1 \leq j \leq d} Q_{X^j}$ does not hold, as shown by considering X^1, X^2 independent Bernoulli 1/2 variables, and $\theta = (1/\sqrt{2}, 1/\sqrt{2})$: then $Q_{X^j}(3/8) = 1/2$ but $Q_{\langle \theta, X \rangle}(3/8) = 3/4$). While independence yields in general

$$Q_{\langle \theta, X \rangle}(t) \leq \min_{1 \leq j \leq d} Q_{\theta^j X^j}(t) = \min_{1 \leq j \leq d} Q_{X^j}(t/|\theta^j|) \leq \max_{1 \leq j \leq d} Q_{X^j}(\sqrt{d} \cdot t),$$

this bound exhibits an undesirable dependence on the dimension d .

Another way of expressing the “non-atomicity” of the distributions of coordinates X^j , which is stable through linear combinations of independent variables, is the rate of decay of their Fourier transform. Indeed, if the distribution of X^j is atomic, then its characteristic function will not vanish at infinity. Proposition 6 below, which relies on an inequality by Esséen, provides uniform anti-concentration for one-dimensional marginals $\langle \theta, X \rangle$ in terms of the Fourier transform of the X^j , which establishes Assumption 1 even without bounded densities. In what follows, we denote Φ_Z the characteristic function of a real random variable Z , defined by $\Phi_Z(\xi) = \mathbb{E}[e^{i\xi Z}]$ for $\xi \in \mathbf{R}$.

Proposition 6. Assume that X^1, \dots, X^d are independent and that there exists constants $C_0 > 0$ and $\alpha \in (0, 1)$ such that, for every $1 \leq j \leq d$ and $\xi \in \mathbf{R}$,

$$|\Phi_{X^j}(\xi)| \leq \frac{1}{(1 + |\xi|/C_0)^\alpha}. \quad (39)$$

Then, $X = (X^1, \dots, X^d)$ satisfies Assumption 1 with $C = 2^{1/\alpha}(2\pi)^{1/\alpha-1}(1-\alpha)^{-1/\alpha}C_0$.

The proof relies on the following lemma.

Lemma 4. Let X^1, \dots, X^d be independent real random variables. Assume that there exists a sub-additive function $g : \mathbf{R}^+ \rightarrow \mathbf{R}$ such that, for every $j = 1, \dots, d$ and $\xi \in \mathbf{R}$,

$$|\Phi_{X^j}(\xi)| \leq \exp(-g(\xi^2)).$$

Then, for every $t \in \mathbf{R}$,

$$Q_X(t) \leq t \cdot \int_{-2\pi/t}^{2\pi/t} \exp(-g(\xi^2)) d\xi. \quad (40)$$

Proof of Lemma 4. For every $\theta \in S^{d-1}$ and $\xi \in \mathbf{R}$, we have, by independence of the X^j ,

$$\begin{aligned} |\Phi_{\langle \theta, X \rangle}(\xi)| &= |\mathbb{E}[e^{i\xi(\theta_1 X^1 + \dots + \theta_d X^d)}]| = |\mathbb{E}[e^{i\xi\theta_1 X^1}]| \dots |\mathbb{E}[e^{i\xi\theta_d X^d}]| \\ &\leq \exp[-(g(\theta_1^2 \xi^2) + \dots + g(\theta_d^2 \xi^2))] \leq \exp(-g(\xi^2)), \end{aligned}$$

where the last inequality uses the sub-additivity of g and the fact that $\theta_1^2 + \dots + \theta_d^2 = \|\theta\|^2 = 1$. Lemma 4 then follows from Esséen's inequality [Ess66], which states that for any real random variable Z ,

$$Q_Z(t) \leq t \cdot \int_{-2\pi/t}^{2\pi/t} |\Phi_Z(\xi)| d\xi. \quad \square$$

Proof of Proposition 6. The functions $g_1 : u \mapsto \alpha \log(1 + u)$ and $g_2 : u \mapsto C_0^{-1} \sqrt{u}$ are concave functions on \mathbf{R}^+ taking the value 0 at 0, and therefore sub-additive. Since g_1 is also increasing, the function $g : u \mapsto g_1 \circ g_2(u) = \alpha \log(1 + C_0^{-1} \sqrt{u})$ is also sub-additive. Condition (39) simply writes $|\Phi_{X^j}(\xi)| \leq \exp(-g(\xi^2))$, so that by Lemma 4

$$Q_X(t) \leq t \int_{-2\pi/t}^{2\pi/t} \frac{1}{(1 + |\xi|/C_0)^\alpha} d\xi \leq 2t \int_0^{2\pi/t} \frac{d\xi}{(\xi/C_0)^\alpha} = 2tC_0^\alpha \left(\frac{2\pi}{t}\right)^{1-\alpha} / (1-\alpha),$$

which implies that $Q_X(t) \leq (Ct)^\alpha$, concluding the proof. \square

4 Proofs from Section 2

In this section, we gather the remaining proofs of results from Section 2 on least squares regression, namely those of Proposition 1, Theorem 1, Proposition 2, Theorem 3 and Proposition 3.

4.1 Preliminary: risk of Ridge and OLS estimators

We start with general expressions for the risk, which will be used several times in the proofs. Here, we assume that (X, Y) is as in Section 2, namely $\mathbb{E}[Y^2] < +\infty$, $\mathbb{E}[\|X\|^2] < +\infty$ and $\Sigma := \mathbb{E}[XX^\top]$ is invertible. Letting $\varepsilon := Y - \langle \beta^*, X \rangle$ denote the error, where $\beta^* := \Sigma^{-1} \mathbb{E}[YX]$ is the risk minimizer, we let $m(X) := \mathbb{E}[\varepsilon|X] = \mathbb{E}[Y|X] - \langle \beta^*, X \rangle$ denote the *misspecification* (or *approximation*) error of the linear model, and $\sigma^2(X) := \text{Var}(\varepsilon|X) = \text{Var}(Y|X)$ denote the conditional variance of the noise.

Lemma 5 (Risk of the Ridge estimator). *Assume that (X, Y) is of the previous form. Let $\lambda \geq 0$, and assume that either $\lambda > 0$ or that P_X is non-degenerate and $n \geq d$. The risk of the Ridge estimator $\hat{\beta}_{\lambda, n}$, defined by*

$$\hat{\beta}_{\lambda, n} := \arg \min_{\beta \in \mathbf{R}^d} \left\{ \frac{1}{n} \sum_{i=1}^n (Y_i - \langle \beta, X_i \rangle)^2 + \lambda \|\beta\|^2 \right\} = (\hat{\Sigma}_n + \lambda I_d)^{-1} \cdot \frac{1}{n} \sum_{i=1}^n Y_i X_i, \quad (41)$$

equals

$$\begin{aligned} \mathbb{E}[\mathcal{E}(\hat{\beta}_{\lambda, n})] &= \mathbb{E} \left[\left\| \frac{1}{n} \sum_{i=1}^n m(X_i) X_i - \lambda \beta^* \right\|_{(\hat{\Sigma}_n + \lambda I_d)^{-1} \Sigma (\hat{\Sigma}_n + \lambda I_d)^{-1}}^2 \right] + \\ &\quad + \frac{1}{n^2} \mathbb{E} \left[\sum_{i=1}^n \sigma^2(X_i) \|X_i\|_{(\hat{\Sigma}_n + \lambda I_d)^{-1} \Sigma (\hat{\Sigma}_n + \lambda I_d)^{-1}}^2 \right]. \end{aligned} \quad (42)$$

Proof. Since $Y_i = \langle \beta^*, X_i \rangle + \varepsilon_i$ for $i = 1, \dots, n$, and since $\langle \beta^*, X_i \rangle X_i = X_i X_i^\top \beta^*$, we have

$$\frac{1}{n} \sum_{i=1}^n Y_i X_i = \hat{\Sigma}_n \beta^* + \frac{1}{n} \sum_{i=1}^n \varepsilon_i X_i. \quad (43)$$

As a result, the excess risk of $\hat{\beta}_{\lambda, n}$ (which is well-defined by the assumptions) writes

$$\begin{aligned} \mathbb{E}[\mathcal{E}(\hat{\beta}_{\lambda, n})] &= \mathbb{E} \left[\left\| (\hat{\Sigma}_n + \lambda I_d)^{-1} \left(\hat{\Sigma}_n \beta^* + \frac{1}{n} \sum_{i=1}^n \varepsilon_i X_i \right) - \beta^* \right\|_{\Sigma}^2 \right] \\ &= \mathbb{E} \left[\left\| (\hat{\Sigma}_n + \lambda I_d)^{-1} \cdot \frac{1}{n} \sum_{i=1}^n \varepsilon_i X_i - \lambda (\hat{\Sigma}_n + \lambda I_d)^{-1} \beta^* \right\|_{\Sigma}^2 \right] \\ &= \mathbb{E} \left[\mathbb{E} \left[\left\| \frac{1}{n} \sum_{i=1}^n \varepsilon_i X_i - \lambda \beta^* \right\|_{(\hat{\Sigma}_n + \lambda I_d)^{-1} \Sigma (\hat{\Sigma}_n + \lambda I_d)^{-1}}^2 \middle| X_1, \dots, X_n \right] \right] \\ &= \mathbb{E} \left[\left\| \frac{1}{n} \sum_{i=1}^n m(X_i) X_i - \lambda \beta^* \right\|_{(\hat{\Sigma}_n + \lambda I_d)^{-1} \Sigma (\hat{\Sigma}_n + \lambda I_d)^{-1}}^2 \right] + \\ &\quad + \frac{1}{n^2} \mathbb{E} \left[\sum_{i=1}^n \sigma^2(X_i) \|X_i\|_{(\hat{\Sigma}_n + \lambda I_d)^{-1} \Sigma (\hat{\Sigma}_n + \lambda I_d)^{-1}}^2 \right] \end{aligned} \quad (44)$$

where (44) is obtained by expanding and using the fact that, for $1 \leq i, j \leq n$ with $i \neq j$,

$$\begin{aligned} \mathbb{E}[\varepsilon_i \varepsilon_j | X_1, \dots, X_n] &= m(X_i) m(X_j), \\ \mathbb{E}[\varepsilon_i^2 | X_1, \dots, X_n] &= m(X_i)^2 + \sigma^2(X_i). \end{aligned} \quad \square$$

In the special case where $\lambda = 0$, the previous risk decomposition becomes:

Lemma 6 (Risk of the OLS estimator). *Assume that P_X is non-degenerate and $n \geq d$. Then,*

$$\mathbb{E}[\mathcal{E}(\hat{\beta}_n^{\text{LS}})] = \mathbb{E} \left[\left\| \frac{1}{n} \sum_{i=1}^n m(X_i) \tilde{X}_i \right\|_{\tilde{\Sigma}_n^{-2}}^2 \right] + \frac{1}{n^2} \mathbb{E} \left[\sum_{i=1}^n \sigma^2(X_i) \|\tilde{X}_i\|_{\tilde{\Sigma}_n^{-2}}^2 \right], \quad (45)$$

where we let $\tilde{X}_i = \Sigma^{-1/2} X_i$ and $\tilde{\Sigma}_n = \Sigma^{-1/2} \hat{\Sigma}_n \Sigma^{-1/2}$.

Proof. This follows from Lemma 5 and the fact that, when $\lambda = 0$, for every $x \in \mathbf{R}^d$,

$$\|x\|_{(\hat{\Sigma}_n + \lambda I_d)^{-1} \Sigma (\hat{\Sigma}_n + \lambda I_d)^{-1}} = \|\Sigma^{-1/2} x\|_{\Sigma^{1/2} \hat{\Sigma}_n^{-1} \Sigma \Sigma^{1/2}} = \|\Sigma^{-1/2} x\|_{\tilde{\Sigma}_n^{-2}}. \quad \square$$

4.2 Proof of Theorem 1 and Proposition 1

Upper bound on the minimax risk. First, we provide an upper bound on the maximum risk of the least-squares estimator over the class $\mathcal{P}_{\text{well}}(P_X, \sigma^2)$. As in Theorem 1, we assume that $n \geq d$ and that P_X is non-degenerate. Let $(X, Y) \sim P \in \mathcal{P}_{\text{well}}(P_X, \sigma^2)$, so that $m(X) = 0$ and $\sigma^2(X) \leq \sigma^2$. It then follows from Lemma 6 that

$$\mathbb{E}[\mathcal{E}(\hat{\beta}_n^{\text{LS}})] \leq \frac{\sigma^2}{n^2} \mathbb{E} \left[\sum_{i=1}^n \sigma^2(X_i) \|\tilde{X}_i\|_{\tilde{\Sigma}_n^{-2}}^2 \right] = \frac{\sigma^2}{n^2} \mathbb{E} \left[\text{Tr} \left(\tilde{\Sigma}_n^{-2} \sum_{i=1}^n \tilde{X}_i \tilde{X}_i^\top \right) \right] = \frac{\sigma^2}{n} \text{Tr}(\tilde{\Sigma}_n^{-1}).$$

Hence, the maximum excess risk of the OLS estimator $\hat{\beta}_n^{\text{LS}}$ over the class $\mathcal{P}_{\text{well}}(P_X, \sigma^2)$ (and consequently, the minimax risk over this class) is at most $\sigma^2 \mathbb{E}[\text{Tr}(\tilde{\Sigma}_n^{-1})]/n$.

Lower bound on the minimax risk. We will now provide a lower bound on the minimax risk over $\mathcal{P}_{\text{Gauss}}(P_X, \sigma^2)$. We will in fact establish the lower bound both in the setting of Theorem 1 (namely, P_X is non-degenerate and $n \geq d$) and that of Proposition 1 (the remaining cases). In particular, we do not assume for now that P_X is non-degenerate or that $n \geq d$.

For $\beta^* \in \mathbf{R}^d$, let P_{β^*} denote the joint distribution of (X, Y) where $X \sim P_X$ and $Y = \langle \beta^*, X \rangle + \varepsilon$ with $\varepsilon \sim \mathcal{N}(0, \sigma^2)$ independent of X . Now, consider the decision problem with model $\mathcal{P}_{\text{Gauss}}(P_X, \sigma^2) = \{P_{\beta^*} : \beta^* \in \mathbf{R}^d\}$, decision space \mathbf{R}^d and loss function $\mathcal{L}(\beta^*, \beta) = \mathcal{E}_{P_{\beta^*}}(\beta) = \|\beta - \beta^*\|_{\Sigma}^2$. Let $\mathcal{R}(\beta^*, \hat{\beta}_n) = \mathbb{E}_{\beta^*}[\mathcal{L}(\beta^*, \hat{\beta}_n)]$ denote the risk under P_{β^*} of a decision rule $\hat{\beta}_n$ (that is, an estimator of β^* using an i.i.d. sample of size n from P_{β^*}), namely its expected excess risk. Consider the prior $\Pi_\lambda = \mathcal{N}(0, \sigma^2/(\lambda n)I_d)$ on $\mathcal{P}_{\text{Gauss}}(P_X, \sigma^2)$. A standard computation (see, e.g., [GCS⁺13]) shows that the posterior $\Pi_\lambda(\cdot | (X_1, Y_1), \dots, (X_n, Y_n))$ is $\mathcal{N}(\hat{\beta}_{\lambda, n}, (\sigma^2/n) \cdot (\hat{\Sigma}_n + \lambda I_d)^{-1})$. Since the loss function \mathcal{L} is quadratic, the Bayes estimator under Π_λ is the expectation of the posterior, which is $\hat{\beta}_{\lambda, n}$. Hence, using the comparison between minimax and Bayes risks:

$$\inf_{\hat{\beta}_n} \sup_{P_{\beta^*} \in \mathcal{P}_{\text{Gauss}}(P_X, \sigma^2)} \mathcal{R}(\beta^*, \hat{\beta}_n) \geq \inf_{\hat{\beta}_n} \mathbb{E}_{\beta^* \sim \Pi_\lambda} [\mathcal{R}(\beta^*, \hat{\beta}_n)] = \mathbb{E}_{\beta^* \sim \Pi_\lambda} [\mathcal{R}(\beta^*, \hat{\beta}_{\lambda, n})], \quad (46)$$

where the infimum is over all estimators $\hat{\beta}_n$. Note that the left-hand side of (46) is simply the minimax excess risk over $\mathcal{P}_{\text{Gauss}}(P_X, \sigma^2)$. On the other hand, applying Lemma 5 with $m(X) = 0$ and $\sigma^2(X) = \sigma^2$ and noting that

$$\begin{aligned} \mathbb{E} \left[\sum_{i=1}^n \|X_i\|_{(\hat{\Sigma}_n + \lambda I_d)^{-1} \Sigma (\hat{\Sigma}_n + \lambda I_d)^{-1}}^2 \right] &= \mathbb{E} \left[\text{Tr} \left\{ (\hat{\Sigma}_n + \lambda I_d)^{-1} \Sigma (\hat{\Sigma}_n + \lambda I_d)^{-1} \sum_{i=1}^n X_i X_i^\top \right\} \right] \\ &= n \mathbb{E} [\text{Tr} \{ (\hat{\Sigma}_n + \lambda I_d)^{-1} \Sigma (\hat{\Sigma}_n + \lambda I_d)^{-1} \hat{\Sigma}_n \}], \end{aligned}$$

we obtain

$$\mathcal{R}(\beta^*, \hat{\beta}_{\lambda, n}) = \lambda^2 \mathbb{E} \left[\|\beta^*\|_{(\hat{\Sigma}_n + \lambda I_d)^{-1} \Sigma (\hat{\Sigma}_n + \lambda I_d)^{-1}}^2 \right] + \frac{\sigma^2}{n} \mathbb{E} [\text{Tr} \{ (\hat{\Sigma}_n + \lambda I_d)^{-1} \Sigma (\hat{\Sigma}_n + \lambda I_d)^{-1} \hat{\Sigma}_n \}].$$

This implies that

$$\begin{aligned} \mathbb{E}_{\beta^* \sim \Pi_\lambda} [\mathcal{R}(\beta^*, \hat{\beta}_{\lambda, n})] &= \mathbb{E}_{\beta^* \sim \Pi_\lambda} \left[\lambda^2 \mathbb{E} \left[\|\beta^*\|_{(\hat{\Sigma}_n + \lambda I_d)^{-1} \Sigma (\hat{\Sigma}_n + \lambda I_d)^{-1}}^2 \right] \right] + \\ &\quad + \frac{\sigma^2}{n} \mathbb{E} [\text{Tr} \{ (\hat{\Sigma}_n + \lambda I_d)^{-1} \Sigma (\hat{\Sigma}_n + \lambda I_d)^{-1} \hat{\Sigma}_n \}] \end{aligned} \quad (47)$$

where \mathbb{E} simply denotes the expectation with respect to $(X_1, \dots, X_n) \sim P_X^n$. Now, by Fubini's theorem, and since $\mathbb{E}_{\beta^* \sim \Pi_\lambda}[\beta^*(\beta^*)^\top] = \sigma^2/(\lambda n)I_d$, we have

$$\begin{aligned} & \mathbb{E}_{\beta^* \sim \Pi_\lambda} \left[\lambda^2 \mathbb{E} \left[\|\beta^*\|_{(\hat{\Sigma}_n + \lambda I_d)^{-1} \Sigma (\hat{\Sigma}_n + \lambda I_d)^{-1}}^2 \right] \right] \\ &= \lambda^2 \cdot \mathbb{E} \left[\mathbb{E}_{\beta^* \sim \Pi_\lambda} \left[\text{Tr} \{ (\hat{\Sigma}_n + \lambda I_d)^{-1} \Sigma (\hat{\Sigma}_n + \lambda I_d)^{-1} \beta^* (\beta^*)^\top \} \right] \right] \\ &= \frac{\sigma^2}{n} \mathbb{E} [\text{Tr} \{ (\hat{\Sigma}_n + \lambda I_d)^{-1} \Sigma (\hat{\Sigma}_n + \lambda I_d)^{-1} \lambda I_d \}]. \end{aligned} \quad (48)$$

Plugging (48) into (47) shows that the Bayes risk under Π_λ equals

$$\frac{\sigma^2}{n} \mathbb{E} [\text{Tr} \{ (\hat{\Sigma}_n + \lambda I_d)^{-1} \Sigma (\hat{\Sigma}_n + \lambda I_d)^{-1} (\hat{\Sigma}_n + \lambda I_d) \}] = \frac{\sigma^2}{n} \mathbb{E} [\text{Tr} \{ (\hat{\Sigma}_n + \lambda I_d)^{-1} \Sigma \}]. \quad (49)$$

Hence, by (46) the minimax risk is larger than $(\sigma^2/n) \cdot \mathbb{E} [\text{Tr} \{ (\hat{\Sigma}_n + \lambda I_d)^{-1} \Sigma \}]$ for every $\lambda > 0$. We will now distinguish the settings of Theorem 1 and Proposition 1.

Degenerate case. First, we assume that P_X is degenerate or that $n < d$. By Definition 1, with probability $p > 0$, the matrix $\hat{\Sigma}_n$ is non-invertible. When this occurs, let $\theta \in \mathbf{R}^d$ be such that $\|\theta\| = 1$ and $\hat{\Sigma}_n(\Sigma^{-1/2}\theta) = 0$. We then have, for every $\lambda > 0$,

$$\langle \Sigma^{-1/2}(\hat{\Sigma}_n + \lambda I_d) \Sigma^{-1/2} \theta, \theta \rangle = 0 + \lambda \|\Sigma^{-1/2} \theta\|^2 \leq \lambda \cdot \lambda_{\min}^{-1},$$

where $\lambda_{\min} = \lambda_{\min}(\Sigma)$ denotes the smallest eigenvalue of Σ . This implies that

$$\text{Tr} \{ \Sigma^{1/2} (\hat{\Sigma}_n + \lambda I_d)^{-1} \Sigma^{1/2} \} \geq \lambda_{\max}(\Sigma^{1/2} (\hat{\Sigma}_n + \lambda I_d)^{-1} \Sigma^{1/2}) = \lambda_{\min}^{-1}(\Sigma^{-1/2} (\hat{\Sigma}_n + \lambda I_d) \Sigma^{-1/2}) \geq \frac{\lambda_{\min}}{\lambda}$$

so that

$$\frac{\sigma^2}{n} \mathbb{E} [\text{Tr} \{ (\hat{\Sigma}_n + \lambda I_d)^{-1} \Sigma \}] \geq \frac{\sigma^2}{n} \cdot p \cdot \frac{\lambda_{\min}}{\lambda}. \quad (50)$$

Recalling that the left-hand side of equation (50) is a lower bound on the minimax risk for every $\lambda > 0$, and noting that the right-hand side tends to $+\infty$ as $\lambda \rightarrow 0$, shows that the minimax risk is infinite as claimed in Proposition 1.

Non-degenerate case. Now, assume that P_X is non-degenerate and that $n \geq d$. By Definition 1, $\hat{\Sigma}_n$ is invertible almost surely. In addition, $\text{Tr} \{ (\hat{\Sigma}_n + \lambda I_d)^{-1} \Sigma \} = \text{Tr} \{ (\Sigma^{-1/2} \hat{\Sigma}_n \Sigma^{-1/2} + \lambda \Sigma^{-1})^{-1} \}$ is decreasing in λ (since $\lambda \mapsto \Sigma^{-1/2} \hat{\Sigma}_n \Sigma^{-1/2} + \lambda \Sigma^{-1}$ is increasing in λ), positive, and converges as $\lambda \rightarrow 0^+$ to $\text{Tr}(\tilde{\Sigma}_n^{-1})$. By the monotone convergence theorem, it follows that

$$\lim_{\lambda \rightarrow 0^+} \frac{\sigma^2}{n} \mathbb{E} [\text{Tr} \{ (\hat{\Sigma}_n + \lambda I_d)^{-1} \Sigma \}] = \frac{\sigma^2}{n} \mathbb{E} [\text{Tr}(\tilde{\Sigma}_n^{-1})], \quad (51)$$

where the limit in the right-hand side belongs to $(0, +\infty]$. Since the left-hand side is a lower bound on the minimax risk, the minimax risk over $\mathcal{P}_{\text{Gauss}}(P_X, \sigma^2)$ is larger than $(\sigma^2/n) \mathbb{E} [\text{Tr}(\tilde{\Sigma}_n^{-1})]$.

Conclusion of the proof. Since $\mathcal{P}_{\text{Gauss}}(P_X, \sigma^2) \subset \mathcal{P}_{\text{well}}(P_X, \sigma^2)$, the minimax risk over $\mathcal{P}_{\text{well}}(P_X, \sigma^2)$ is at least as large as that over $\mathcal{P}_{\text{Gauss}}(P_X, \sigma^2)$. In the case when P_X is degenerate or $n < d$, we showed that the minimax risk over $\mathcal{P}_{\text{Gauss}}(P_X, \sigma^2)$ is infinite, establishing Proposition 1. In the case when P_X is non-degenerate and $n \geq d$, we showed that the minimax risk over $\mathcal{P}_{\text{well}}(P_X, \sigma^2)$ is smaller than $(\sigma^2/n) \mathbb{E} [\text{Tr}(\tilde{\Sigma}_n^{-1})]$ and that the minimax risk over $\mathcal{P}_{\text{Gauss}}(P_X, \sigma^2)$ is larger than the same quantity, so that both minimax risks agree and are equal to $(\sigma^2/n) \mathbb{E} [\text{Tr}(\tilde{\Sigma}_n^{-1})]$, as claimed in Theorem 1.

4.3 Proof of Theorem 3

The proof starts with the following lemma.

Lemma 7. *For any positive symmetric $d \times d$ matrix A and $p \in [1, 2]$,*

$$\mathrm{Tr}(A^{-1}) + \mathrm{Tr}(A) - 2d \leq \max(1, \lambda_{\min}(A)^{-1}) \cdot \mathrm{Tr}(|A - I_d|^{2/p}). \quad (52)$$

Proof of Lemma 7. Let us start by showing that, for every $a > 0$,

$$a^{-1} + a - 2 \leq \max(1, a^{-1}) \cdot |a - 1|^{2/p}. \quad (53)$$

Multiplying both sides of (53) by $a > 0$, it amounts to

$$(a - 1)^2 = 1 + a^2 - 2a \leq \max(a, 1) \cdot |a - 1|^{2/p},$$

namely to $|a - 1|^{2-2/p} \leq \max(a, 1)$. For $a \in (0, 2]$, this inequality holds since $|a - 1| \leq 1$ and $2 - 2/p \geq 0$, so that $|a - 1|^{2-2/p} \leq 1 \leq \max(a, 1)$. For $a \geq 2$, the inequalities $|a - 1| \geq 2$ and $2 - 2/p \leq 1$ imply that $|a - 1|^{2-2/p} \leq |a - 1| \leq a \leq \max(a, 1)$. This establishes (53).

Now, let $a_1, \dots, a_d > 0$ be the eigenvalues of A . Without loss of generality, assume that $a_d = \min_j(a_j) = \lambda_{\min}(A)$. Then, by inequality (53) and the bound $\max(1, a_j^{-1}) \leq \max(1, a_d^{-1})$, we have

$$\mathrm{Tr}(A^{-1}) + \mathrm{Tr}(A) - 2d = \sum_{j=1}^d (a_j^{-1} + a_j - 2) \leq \max(1, a_d^{-1}) \sum_{j=1}^d |a_j - 1|^{2/p},$$

which is precisely the desired inequality (52). \square

Proof of Theorem 3. Let $p \in (1, 2]$ which will be determined later, and denote $q := p/(p - 1)$ its complement. Applying Lemma 7 to $A = \tilde{\Sigma}_n$ yields:

$$\mathrm{Tr}(\tilde{\Sigma}_n^{-1}) + \mathrm{Tr}(\tilde{\Sigma}_n) - 2d \leq \max(1, \lambda_{\min}(\tilde{\Sigma}_n)^{-1}) \cdot \mathrm{Tr}(|\tilde{\Sigma}_n - I_d|^{2/p}).$$

Also, note that both random variables are non-negative, since $\mathrm{Tr}(\tilde{\Sigma}_n^{-1}) + \mathrm{Tr}(\tilde{\Sigma}_n) - 2d = \mathrm{Tr}[(\tilde{\Sigma}_n^{1/2} - \tilde{\Sigma}_n^{-1/2})^2] \geq 0$. Hence, their expectations are well-defined in $[0, +\infty]$, and since $\mathbb{E}[\mathrm{Tr}(\tilde{\Sigma}_n)] = d$, taking the expectation in the above bound and dividing by d yields:

$$\begin{aligned} \frac{1}{d} \cdot \mathbb{E}[\mathrm{Tr}(\tilde{\Sigma}_n^{-1})] - 1 &\leq \mathbb{E}\left[\max(1, \lambda_{\min}(\tilde{\Sigma}_n)^{-1}) \cdot \frac{1}{d} \mathrm{Tr}(|\tilde{\Sigma}_n - I_d|^{2/p})\right] \\ &\leq \mathbb{E}\left[\max(1, \lambda_{\min}(\tilde{\Sigma}_n)^{-1})^q\right]^{1/q} \cdot \mathbb{E}\left[\left(\frac{1}{d} \mathrm{Tr}(|\tilde{\Sigma}_n - I_d|^{2/p})\right)^p\right]^{1/p} \end{aligned} \quad (54)$$

$$\leq \mathbb{E}\left[\max(1, \lambda_{\min}(\tilde{\Sigma}_n)^{-q})\right]^{1/q} \cdot \mathbb{E}\left[\frac{1}{d} \mathrm{Tr}((\tilde{\Sigma}_n - I_d)^2)\right]^{1/p} \quad (55)$$

where (54) comes from Hölder's inequality, while (55) is obtained by noting that $x \mapsto x^p$ is convex and that $(1/d)\mathrm{Tr}(A)$ is the average of the eigenvalues of the symmetric matrix A . Next,

$$\begin{aligned} \mathbb{E}\left[\frac{1}{d} \mathrm{Tr}((\tilde{\Sigma}_n - I_d)^2)\right] &= \frac{1}{d} \mathrm{Tr}\left\{\mathbb{E}\left[\left(\frac{1}{n} \sum_{i=1}^n (\tilde{X}_i \tilde{X}_i^\top - I_d)\right)^2\right]\right\} \\ &= \frac{1}{n^2 d} \mathrm{Tr}\left\{\sum_{1 \leq i, j \leq n} \mathbb{E}[(\tilde{X}_i \tilde{X}_i^\top - I_d)(\tilde{X}_j \tilde{X}_j^\top - I_d)]\right\} \\ &= \frac{1}{nd} \mathrm{Tr}\left\{\mathbb{E}[(\tilde{X} \tilde{X}^\top - I_d)^2]\right\}, \end{aligned} \quad (56)$$

where we used in (56) the fact that, for $i \neq j$, $\mathbb{E}[(\tilde{X}_i \tilde{X}_i^\top - I_d)(\tilde{X}_j \tilde{X}_j^\top - I_d)] = \mathbb{E}[\tilde{X}_i \tilde{X}_i^\top - I_d] \mathbb{E}[\tilde{X}_j \tilde{X}_j^\top - I_d] = 0$. Now, observe that for $x \in \mathbf{R}^d$, $xx^\top - I_d$ has eigenvalue $\|x\|^2 - 1$ in the direction of x , and eigenvalue -1 in any orthogonal direction. It follows that

$$\text{Tr}\{(xx^\top - I_d)^2\} = (\|x\|^2 - 1)^2 + (d-1) \cdot (-1)^2 = \|x\|^4 - 2\|x\|^2 + d,$$

so that (56) becomes, recalling that $\mathbb{E}[\|\tilde{X}\|^2] = d$ and $\mathbb{E}[\|\tilde{X}\|^4] \leq \kappa d^2$ (Assumption 2),

$$\mathbb{E}\left[\frac{1}{d} \text{Tr}((\tilde{\Sigma}_n - I_d)^2)\right] = \frac{1}{nd} \left(\mathbb{E}[\|\tilde{X}\|^4] - 2\mathbb{E}[\|\tilde{X}\|^2] + d \right) = \frac{1}{n} \left(\frac{1}{d} \mathbb{E}[\|\tilde{X}\|^4] - 1 \right) \leq \frac{\kappa d}{n}. \quad (57)$$

On the other, recall that \tilde{X} satisfies Assumption 1 and that $n \geq \max(6d/\alpha, 12/\alpha)$. Hence, letting $C' \geq 1$ be the constant in Theorem 4, we have by Corollary 4:

$$\mathbb{E}[\max(1, \lambda_{\min}(\tilde{\Sigma}_n)^{-q})] \leq 2C'^q. \quad (58)$$

Finally, plugging the bounds (57) and (58) into (55) and recalling that $1/p = 1 - 1/q = 1 - 2/(\alpha'n)$, we obtain

$$\frac{1}{d} \cdot \mathbb{E}[\text{Tr}(\tilde{\Sigma}_n^{-1})] - 1 \leq (2C'^q)^{1/q} \cdot \left(\frac{\kappa d}{n}\right)^{1/p} \leq 2C' \cdot \frac{\kappa d}{n} \cdot \left(\frac{n}{\kappa d}\right)^{2/(\alpha'n)}. \quad (59)$$

Now, since $\kappa = \mathbb{E}[\|\tilde{X}\|^4]/\mathbb{E}[\|\tilde{X}\|^2]^2 \geq 1$ and $d \geq 1$,

$$\left(\frac{n}{\kappa d}\right)^{2/(\alpha'n)} \leq n^{2/(\alpha'n)} = \exp\left(\frac{2 \log n}{\alpha'n}\right).$$

An elementary analysis shows that the function $g : x \mapsto \log x/x$ is increasing on $(0, e]$ and decreasing on $[e, +\infty)$. Hence, if $x, y > 1$ satisfy $x \geq y \log y \geq e$, then

$$\frac{\log x}{x} \leq \frac{\log y + \log \log y}{y \log y} \leq \frac{1 + e^{-1}}{y}$$

where we used $\log \log y / \log y \leq g(e) = e^{-1}$. Here by assumption $n \geq 12\alpha^{-1} \log(12\alpha^{-1}) = 2\alpha'^{-1} \log(2\alpha'^{-1})$, and thus $\log n/n \leq (1 + e^{-1})/(2\alpha')$, so that

$$\left(\frac{n}{\kappa d}\right)^{2/(\alpha'n)} \leq \exp\left(\frac{2}{\alpha'} \cdot \frac{1 + e^{-1}}{2\alpha'}\right) = \exp(1 + e^{-1}) \leq 4.$$

Plugging this inequality into (59) yields the desired bound (22). Equation (23) then follows by Theorem 1. \square

4.4 Proof of Proposition 3

Recall that, by Lemma 6, we have

$$\mathbb{E}[\mathcal{E}(\hat{\beta}_n^{\text{LS}})] = \mathbb{E}\left[\left\|\frac{1}{n} \sum_{i=1}^n m(X_i) \Sigma^{-1/2} X_i\right\|_{\tilde{\Sigma}_n^{-2}}^2\right] + \frac{1}{n^2} \mathbb{E}\left[\sum_{i=1}^n \sigma^2(X_i) \|\Sigma^{-1/2} X_i\|_{\tilde{\Sigma}_n^{-2}}^2\right]. \quad (60)$$

Now, since $\tilde{\Sigma}_n^{-2} \leq \lambda_{\min}(\tilde{\Sigma}_n)^{-2} I_d$, we have for every random variable V_n :

$$\begin{aligned} \mathbb{E}[\|V_n\|_{\tilde{\Sigma}_n^{-2}}^2] &\leq \mathbb{E}[\|V_n\|^2] + \mathbb{E}[\{\lambda_{\min}(\tilde{\Sigma}_n)^{-2} - 1\}_+ \cdot \|V_n\|^2] \\ &\leq \mathbb{E}[\|V_n\|^2] + \mathbb{E}[\{\lambda_{\min}(\tilde{\Sigma}_n)^{-2} - 1\}_+^2]^{1/2} \cdot \mathbb{E}[\|V_n\|^4]^{1/2}, \end{aligned} \quad (61)$$

where (61) follows from the Cauchy-Schwarz inequality. Letting $V_n = \sigma(X_i)\Sigma^{-1/2}X_i$, we obtain from (61)

$$\begin{aligned} & \frac{1}{n^2} \mathbb{E} \left[\sum_{i=1}^n \sigma^2(X_i) \|\Sigma^{-1/2}X_i\|_{\tilde{\Sigma}_n^{-2}}^2 \right] \\ & \leq \frac{1}{n} \mathbb{E} [\sigma^2(X) \|\Sigma^{-1/2}X\|^2] + \frac{1}{n} \mathbb{E} [\{\lambda_{\min}(\tilde{\Sigma}_n)^{-2} - 1\}_+^2]^{1/2} \mathbb{E} [\sigma^4(X) \|\Sigma^{-1/2}X\|^4]^{1/2}. \end{aligned} \quad (62)$$

On the other hand, let $V_n = n^{-1} \sum_{i=1}^n m(X_i) \Sigma^{-1/2} X_i$; we have, since $\mathbb{E}[m(X_i)X_i] = \mathbb{E}[\varepsilon_i X_i] = 0$,

$$\begin{aligned} \mathbb{E}[\|V_n\|^2] &= \mathbb{E} \left[\left\| \frac{1}{n} \sum_{i=1}^n m(X_i) X_i \right\|_{\Sigma^{-1}}^2 \right] \\ &= \frac{1}{n^2} \sum_{1 \leq i, j \leq n} \mathbb{E} [\langle m(X_i) X_i, m(X_j) X_j \rangle_{\Sigma^{-1}}] \\ &= \frac{1}{n^2} \sum_{i=1}^n \mathbb{E} [m(X_i)^2 \|\Sigma^{-1/2} X_i\|^2] + \frac{1}{n^2} \sum_{i \neq j} \langle \mathbb{E}[m(X_i) X_i], \mathbb{E}[m(X_j) X_j] \rangle_{\Sigma^{-1}} \\ &= \frac{1}{n} \mathbb{E} [m(X)^2 \|\Sigma^{-1/2} X\|^2]. \end{aligned} \quad (63)$$

In addition,

$$\mathbb{E}[\|V_n\|^4] = \frac{1}{n^4} \sum_{1 \leq i, j, k, l \leq n} \mathbb{E} [\langle m(X_i) X_i, m(X_j) X_j \rangle_{\Sigma^{-1}} \langle m(X_k) X_k, m(X_l) X_l \rangle_{\Sigma^{-1}}].$$

Now, by independence and since $\mathbb{E}[m(X)X] = 0$, each term in the sum above where one index among i, j, k, l is distinct from the others cancels. We therefore have

$$\begin{aligned} \mathbb{E}[\|V_n\|^4] &= \frac{1}{n^4} \sum_{i=1}^n \mathbb{E} [\|m(X_i) X_i\|_{\Sigma^{-1}}^4] + \frac{2}{n^4} \sum_{1 \leq i < j \leq n} \mathbb{E} [\|m(X_i) X_i\|_{\Sigma^{-1}}^2 \|m(X_j) X_j\|_{\Sigma^{-1}}^2] + \\ &+ \frac{4}{n^4} \sum_{1 \leq i < j \leq n} \mathbb{E} [\langle m(X_i) X_i, m(X_j) X_j \rangle_{\Sigma^{-1}}^2] \\ &\leq \frac{1}{n^4} \sum_{i=1}^n \mathbb{E} [\|m(X_i) X_i\|_{\Sigma^{-1}}^4] + \frac{6}{n^4} \sum_{1 \leq i < j \leq n} \mathbb{E} [\|m(X_i) X_i\|_{\Sigma^{-1}}^2 \|m(X_j) X_j\|_{\Sigma^{-1}}^2] \quad (64) \\ &= \frac{1}{n^3} \cdot \mathbb{E} [m(X)^4 \|\Sigma^{-1/2} X\|^4] + \frac{6}{n^4} \cdot \frac{n(n-1)}{2} \cdot \mathbb{E} [m(X)^2 \|X\|_{\Sigma^{-1}}^2]^2 \\ &\leq \frac{1}{n^3} \cdot \mathbb{E} [m(X)^4 \|\Sigma^{-1/2} X\|^4] + \frac{3}{n^2} \cdot \mathbb{E} [m(X)^2 \|\Sigma^{-1/2} X\|^2]^2 \\ &\leq \frac{4}{n^2} \cdot \mathbb{E} [m(X)^4 \|\Sigma^{-1/2} X\|^4] \end{aligned} \quad (65)$$

where (64) and (65) rely on the Cauchy-Schwarz inequality. Hence, it follows from (61), (63) and (65) that

$$\begin{aligned} & \mathbb{E} \left[\left\| \frac{1}{n} \sum_{i=1}^n m(X_i) \Sigma^{-1/2} X_i \right\|_{\tilde{\Sigma}_n^{-2}}^2 \right] \\ & \leq \frac{1}{n} \mathbb{E} [m(X)^2 \|\Sigma^{-1/2} X\|^2] + \mathbb{E} [\{\lambda_{\min}(\tilde{\Sigma}_n)^{-2} - 1\}_+^2]^{1/2} \cdot \left(\frac{4}{n^2} \cdot \mathbb{E} [m(X)^4 \|\Sigma^{-1/2} X\|^4] \right)^{1/2} \\ & \leq \frac{1}{n} \mathbb{E} [m(X)^2 \|\Sigma^{-1/2} X\|^2] + \frac{2}{n} \mathbb{E} [\{\lambda_{\min}(\tilde{\Sigma}_n)^{-2} - 1\}_+^2]^{1/2} \mathbb{E} [m(X)^4 \|\Sigma^{-1/2} X\|^4]^{1/2}. \end{aligned} \quad (66)$$

Plugging (62) and (66) into the decomposition (60) yields:

$$\begin{aligned} \mathbb{E}[\mathcal{E}(\hat{\beta}_n^{\text{LS}})] &\leq \frac{1}{n} \mathbb{E}[(m(X)^2 + \sigma^2(X)) \|\Sigma^{-1/2} X\|^2] + \frac{1}{n} \mathbb{E}[\{\lambda_{\min}(\tilde{\Sigma}_n)^{-2} - 1\}_+^2]^{1/2} \times \\ &\quad \times \left(\mathbb{E}[\sigma^4(X) \|\Sigma^{-1/2} X\|^4]^{1/2} + 2 \mathbb{E}[m(X)^4 \|\Sigma^{-1/2} X\|^4]^{1/2} \right) \end{aligned} \quad (67)$$

Oliveira's subgaussian bound. [Oli16] showed that, under Assumption 3, we have

$$\mathbb{P}(\lambda_{\min}(\hat{\Sigma}_n) \geq 1 - \varepsilon) \geq 1 - \delta$$

provided that

$$n \geq \frac{81\kappa(d + 2\log(2/\delta))}{\varepsilon^2}.$$

This can be rewritten as:

$$\mathbb{P}\left(\lambda_{\min}(\hat{\Sigma}_n) < 1 - 9\kappa^{1/2} \sqrt{\frac{(d + 2\log(2/\delta))}{n}}\right) \leq \delta. \quad (68)$$

Bound on the remaining term. Since the function $x \mapsto x^2$ is 2-Lipschitz on $[0, 1]$, we have $(x^{-2} - 1)_+ = (1 - x^2)_+/x^2 \leq 2(1 - x)_+/x^2$ for $x > 0$, so that by Cauchy-Schwarz,

$$\begin{aligned} \mathbb{E}[\{\lambda_{\min}(\hat{\Sigma}_n)^{-2} - 1\}_+^2]^{1/2} &\leq \mathbb{E}\left[\frac{4\{1 - \lambda_{\min}(\hat{\Sigma}_n)\}_+^2}{\lambda_{\min}(\hat{\Sigma}_n)^4}\right]^{1/2} \\ &\leq 2\mathbb{E}[\{1 - \lambda_{\min}(\hat{\Sigma}_n)\}_+^4]^{1/4} \mathbb{E}[\lambda_{\min}(\hat{\Sigma}_n)^{-8}]^{1/4}. \end{aligned} \quad (69)$$

First, note that

$$\begin{aligned} \mathbb{E}[\{1 - \lambda_{\min}(\hat{\Sigma}_n)\}_+^4] &= \int_0^\infty \mathbb{P}(\{1 - \lambda_{\min}(\hat{\Sigma}_n)\}_+^4 \geq u) du \\ &= \int_0^1 \mathbb{P}(\lambda_{\min}(\hat{\Sigma}_n) \leq 1 - u^{1/4}) du \\ &= \int_0^1 \mathbb{P}(\lambda_{\min}(\hat{\Sigma}_n) \leq 1 - v^{1/2}) 2v dv. \end{aligned} \quad (70)$$

Now, let $v^{1/2} = 9\kappa^{1/2} \sqrt{(d + 2\log(2/\delta))/n}$, so that the bound (68) yields $\mathbb{P}(\lambda_{\min}(\hat{\Sigma}_n) \leq 1 - v^{1/2}) \leq \delta$. We have, equivalently,

$$\delta = 2 \exp\left(-\frac{n}{162\kappa}\left(v - \frac{81\kappa d}{n}\right)\right) \leq 2 \exp\left(-\frac{n}{324\kappa}v\right)$$

as long as $v \geq 162\kappa d/n$. Plugging this inequality into (70) yields

$$\begin{aligned} \mathbb{E}[\{1 - \lambda_{\min}(\hat{\Sigma}_n)\}_+^4] &\leq \int_0^{\min(162\kappa d/n, 1)} 2v dv + \int_{\min(162\kappa d/n, 1)}^1 2 \exp\left(-\frac{n}{324\kappa}v\right) 2v dv \\ &\leq \left(\frac{162\kappa d}{n}\right)^2 + \left(\frac{324\kappa}{n}\right)^2 \int_0^\infty 4 \exp(-w) w dw \\ &= \left(\frac{162\kappa d}{n}\right)^2 + 4 \left(\frac{324\kappa}{n}\right)^2 \end{aligned}$$

so that, using the inequality $(x + y)^{1/4} \leq x^{1/4} + y^{1/4}$,

$$\mathbb{E}[\{1 - \lambda_{\min}(\hat{\Sigma}_n)\}_+^4]^{1/4} \leq 9\sqrt{\frac{2\kappa d}{n}} + 18\sqrt{\frac{2\kappa}{n}} \leq 27\sqrt{\frac{2\kappa d}{n}}. \quad (71)$$

Also, by Corollary 4 and the fact that $\alpha n/12 \geq 8$, $\mathbb{E}[\lambda_{\min}(\widehat{\Sigma}_n)^{-8}] \leq 2C'^8$, so that inequality (69) becomes

$$\mathbb{E}[\{\lambda_{\min}(\widehat{\Sigma}_n)^{-2} - 1\}_+^2]^{1/2} \leq 2 \times 27 \sqrt{\frac{2\kappa d}{n}} \times 2^{1/4} C'^2 \leq 92C'^2 \sqrt{\frac{\kappa d}{n}}. \quad (72)$$

Final bound. Now, let $\chi > 0$ as in Proposition 3. Since

$$\mathbb{E}[\varepsilon^2|X] = m(X)^2 + \sigma^2(X) \geq \max(m(X)^2, \sigma^2(X)),$$

we have

$$\max\left(\mathbb{E}[m(X)^4 \|\Sigma^{-1/2} X\|^4], \mathbb{E}[\sigma^4(X) \|\Sigma^{-1/2} X\|^4]\right) \leq \mathbb{E}[\mathbb{E}[\varepsilon^2|X]^2 \|\Sigma^{-1/2} X\|^4] = \chi d^2. \quad (73)$$

Putting the bounds (72) and (73) inside (67) yields

$$\begin{aligned} \mathbb{E}[\mathcal{E}(\widehat{\beta}_n^{\text{LS}})] &\leq \frac{1}{n} \mathbb{E}[(m(X)^2 + \sigma^2(X)) \|\Sigma^{-1/2} X\|^2] + \frac{1}{n} \cdot 92C'^2 \sqrt{\frac{\kappa d}{n}} \cdot 3\sqrt{\chi} d \\ &= \frac{1}{n} \mathbb{E}[(Y - \langle \beta^*, X \rangle)^2 \|\Sigma^{-1/2} X\|^2] + 276C'^2 \sqrt{\kappa \chi} \left(\frac{d}{n}\right)^{3/2}, \end{aligned} \quad (74)$$

where we used the fact that $\mathbb{E}[(Y - \langle \beta^*, X \rangle)^2|X] = m(X)^2 + \sigma^2(X)$. This establishes (24). Finally, if $P \in \mathcal{P}_{\text{mis}}(P_X, \sigma^2)$, then $\mathbb{E}[\varepsilon^2|X] \leq \sigma^2$, so that

$$\chi = \mathbb{E}[\mathbb{E}[\varepsilon^2|X]^2 \|\Sigma^{-1/2} X\|^4]/d^2 \leq \sigma^4 \mathbb{E}[\|\Sigma^{-1/2} X\|^4]/d^2 \leq \sigma^4 \kappa,$$

where we used the fact that $\mathbb{E}[\|\Sigma^{-1/2} X\|^4] \leq \kappa d^2$ by Assumption 3 (see Remark 3). Plugging this inequality, together with $\mathbb{E}[(Y - \langle \beta^*, X \rangle)^2 \|\Sigma^{-1/2} X\|^2] \leq \sigma^2 d$, inside (74), yields the upper bound (25). This concludes the proof.

5 Proof of Theorem 4

5.1 Truncation and small-ball condition

The first step of the proof is to replace X by the truncated vector $X' := (1 \wedge \frac{\sqrt{d}}{\|X\|})X$; likewise, let $X'_i = (1 \wedge \frac{\sqrt{d}}{\|X_i\|})X_i$ for $1 \leq i \leq n$, and $\widehat{\Sigma}'_n := n^{-1} \sum_{i=1}^n X'_i (X'_i)^\top$. Note that $X' (X')^\top \preceq X X^\top$ and $\|X'\| = \sqrt{d} \wedge \|X\|$, so that $\widehat{\Sigma}'_n \preceq \widehat{\Sigma}_n$ and $\mathbb{E}[\|X'\|^2] \leq \mathbb{E}[\|X\|^2] = d$. It follows that $\lambda_{\min}(\widehat{\Sigma}'_n) \leq \lambda_{\min}(\widehat{\Sigma}_n)$, hence it suffices to establish a lower bound for $\lambda_{\min}(\widehat{\Sigma}'_n)$.

In addition, for every $\theta \in S^{d-1}$, $t \in (0, C^{-1})$ and $a \geq 1$,

$$\begin{aligned} \mathbb{P}(|\langle X', \theta \rangle| \leq t) &\leq \mathbb{P}(|\langle X, \theta \rangle| \leq at) + \mathbb{P}\left(\frac{\sqrt{d}}{\|X\|} \leq \frac{1}{a}\right) \\ &\leq (Cat)^\alpha + \mathbb{P}(\|X\| \geq a\sqrt{d}) \\ &\leq (Cat)^\alpha + \frac{\mathbb{E}[\|X\|^2]}{a^2 d} \end{aligned} \quad (75)$$

$$= (Ct)^\alpha a^\alpha + \frac{1}{a^2} \quad (76)$$

where we applied Markov's inequality in (75). In particular, letting $a = (Ct)^{-\alpha/(2+\alpha)}$, inequality (76) becomes

$$\mathbb{P}(|\langle X', \theta \rangle| \leq t) \leq 2(Ct)^{2\alpha/(2+\alpha)}. \quad (77)$$

5.2 Concentration and PAC-Bayesian inequalities

The smallest eigenvalue $\lambda_{\min}(\widehat{\Sigma}'_n)$ of $\widehat{\Sigma}'_n$ may be written as the infimum of an empirical process indexed by the unit sphere $S^{d-1} = \{v \in \mathbf{R}^d : \|v\| = 1\}$:

$$\lambda_{\min}(\widehat{\Sigma}'_n) = \inf_{v \in S^{d-1}} \langle \widehat{\Sigma}'_n v, v \rangle = \inf_{v \in S^{d-1}} \frac{1}{n} \sum_{i=1}^n \langle X'_i, v \rangle^2.$$

Now, recall that the variables $\langle X'_i, \theta \rangle^2$ are i.i.d. and distributed as $\langle X', \theta \rangle^2$ for every $\theta \in S^{d-1}$. The inequality (77) on the left tail of this variable can be expressed in terms of its Laplace transform, through the following lemma:

Lemma 8. *Let Z be a nonnegative random variable. Assume that there exists $\alpha \in (0, 1]$ and $C > 0$ such that, for every $t \geq 0$, $\mathbb{P}(Z \leq t) \leq (Ct)^\alpha$. Then, for every $\lambda > 0$,*

$$\mathbb{E}[\exp(-\lambda Z)] \leq (C/\lambda)^\alpha. \quad (78)$$

Proof of Lemma 8. Since $0 \leq \exp(-\lambda Z) \leq 1$, we have

$$\mathbb{E}[\exp(-\lambda Z)] = \int_0^1 \mathbb{P}(\exp(-\lambda Z) \geq t) dt = \int_0^1 \mathbb{P}\left(Z \leq \frac{\log(1/t)}{\lambda}\right) dt \leq \int_0^1 \left(C \frac{\log(1/t)}{\lambda}\right)^\alpha dt.$$

Now, for $u > 0$, the map $\alpha \mapsto u^\alpha = e^{\alpha \log u}$ is convex on \mathbf{R} , so that $u^\alpha \leq \alpha u + (1 - \alpha)$ for $0 \leq \alpha \leq 1$. It follows that

$$\int_0^1 \log^\alpha(1/t) dt \leq \alpha \int_0^1 (-\log t) dt + (1 - \alpha) = \alpha [-t \log t + t]_0^1 + (1 - \alpha) = 1,$$

which establishes inequality (78). \square

Here, inequality (77) implies that, for every $\theta \in S^{d-1}$,

$$\mathbb{P}(\langle X', \theta \rangle^2 \leq t) = \mathbb{P}(|\langle X', \theta \rangle| \leq \sqrt{t}) \leq 2(C\sqrt{t})^{2\alpha/(2+\alpha)} = 2(C^2 t)^{\alpha/(2+\alpha)}.$$

Hence, Lemma 8 with $Z = \langle X', \theta \rangle^2$ implies that, for every $\lambda > 0$,

$$\mathbb{E}[\exp(-\lambda \langle X', \theta \rangle^2)] \leq 2(C^2/\lambda)^{\alpha/(2+\alpha)}.$$

In other words, for $i = 1, \dots, n$, $\mathbb{E}[\exp(Z_i(\theta))] \leq 1$, where, letting $\alpha' = \alpha/(2 + \alpha)$, we define

$$Z_i(\theta) = -\lambda \langle X'_i, \theta \rangle^2 + \alpha' \log \left(\frac{\lambda}{C^2} \right) - \log 2$$

with $\lambda > 0$ a fixed parameter that will be optimized later. In particular, letting

$$Z(\theta) = Z_1(\theta) + \dots + Z_n(\theta) = n \left[-\lambda \langle \widehat{\Sigma}'_n \theta, \theta \rangle + \alpha' \log \left(\frac{\lambda}{C^2} \right) - \log 2 \right],$$

the independence of $Z_1(\theta), \dots, Z_n(\theta)$ implies that, for every $\theta \in S^{d-1}$,

$$\mathbb{E}[\exp(Z(\theta))] = \mathbb{E}[\exp(Z_1(\theta))] \cdots \mathbb{E}[\exp(Z_n(\theta))] \leq 1. \quad (79)$$

The bound (79) controls the upper tail of $Z(\theta)$ for fixed $\theta \in \Theta$. In order to obtain a uniform control over θ , similarly to [AC11, Oli16] we will use the PAC-Bayesian technique for bounding empirical processes [McA99b, McA99a, Cat07]. For completeness, we include a proof of Lemma 9 (which is a standard bound) below.

Lemma 9 (PAC-Bayesian deviation bound). *Let Θ be a measurable space, and $Z(\theta)$, $\theta \in \Theta$, be a real-valued measurable process. Assume that $\mathbb{E}[\exp Z(\theta)] \leq 1$ for every $\theta \in \Theta$. Let π be a probability distribution on Θ . Then,*

$$\mathbb{P}\left(\forall \rho, \int_{\Theta} Z(\theta) \rho(d\theta) \leq \text{KL}(\rho, \pi) + t\right) \geq 1 - e^{-t}, \quad (80)$$

where ρ spans over all probability distributions on Θ , and $\text{KL}(\rho, \pi) := \int_{\Theta} \log\left(\frac{d\rho}{d\pi}\right) d\rho \in [0, +\infty]$ denotes the Kullback-Leibler divergence between ρ and π , and where we define the integral in (80) to be $-\infty$ when neither the positive nor the negative part are integrable.

Proof of Lemma 9. By integrating the inequality $\mathbb{E}[\exp Z(\theta)] \leq 1$ with respect to π and using the Fubini-Tonelli theorem, we obtain

$$\mathbb{E}\left[\int_{\Theta} \exp Z(\theta) \pi(d\theta)\right] \leq 1. \quad (81)$$

In addition, using the duality between the log-Laplace transform and the Kullback-Leibler divergence (see, e.g., [Cat04, p. 159]):

$$\log \int_{\Theta} \exp(Z(\theta)) \pi(d\theta) = \sup_{\rho} \left\{ \int_{\Theta} Z(\theta) \rho(d\theta) - \text{KL}(\rho, \pi) \right\}$$

where the supremum spans over all probability distributions ρ over Θ , the inequality (81) writes

$$\mathbb{E}\left[\exp \sup_{\rho} \left\{ \int_{\Theta} Z(\theta) \rho(d\theta) - \text{KL}(\rho, \pi) \right\}\right] \leq 1. \quad (82)$$

Applying Markov's inequality to (82) yields the desired bound (80). \square

Here, we let $\Theta = S^{d-1}$ and $Z(\theta)$ as defined above. In addition, we take π to be the uniform distribution on S^{d-1} , and for $v \in S^{d-1}$ and $\gamma > 0$ we define $\Theta(v, \gamma) := \{\theta \in S^{d-1} : \|\theta - v\| \leq \gamma\}$ and let $\rho_{v, \gamma} = \pi(\Theta(v, \gamma))^{-1} \mathbf{1}(\Theta(v, \gamma)) \cdot \pi$ be the uniform distribution over $\Theta(v, \gamma)$. In this case, the PAC-Bayesian bound of Lemma 9 writes: for every $t > 0$, with probability at least $1 - e^{-t}$, for every $v \in S^{d-1}$ and $\gamma > 0$,

$$n \left[-\lambda F_{v, \gamma}(\widehat{\Sigma}'_n) + \alpha' \log\left(\frac{\lambda}{C^2}\right) - \log 2 \right] \leq \text{KL}(\rho_{v, \gamma}, \pi) + t, \quad (83)$$

where we define for every symmetric matrix Σ :

$$F_{v, \gamma}(\Sigma) := \int_{\Theta} \langle \Sigma \theta, \theta \rangle \rho_{v, \gamma}(d\theta). \quad (84)$$

5.3 Control of the approximation term

Now, using the symmetries of the smoothing distributions $\rho_{v, \gamma}$, we will show that, for every $\gamma > 0$, $v \in S^{d-1}$ and symmetric matrix Σ ,

$$F_{v, \gamma}(\Sigma) = (1 - \phi(\gamma)) \langle \Sigma v, v \rangle + \phi(\gamma) \cdot \frac{1}{d} \text{Tr}(\Sigma), \quad (85)$$

where for $\gamma > 0$,

$$\phi(\gamma) := \frac{d}{d-1} \int_{\Theta} (1 - \langle \theta, v \rangle^2) \rho_{v, \gamma}(d\theta) \in [0, d/(d-1)\gamma^2]. \quad (86)$$

First, note that

$$F_{v,\gamma}(\Sigma) = \text{Tr}(\Sigma A_{v,\gamma}), \quad \text{where} \quad A_{v,\gamma} := \int_{\Theta} \theta \theta^\top \rho_{v,\gamma}(\text{d}\theta).$$

In addition, for every isometry $U \in O(d)$ of \mathbf{R}^d and $v \in S^{d-1}$, $\gamma > 0$, the image measure $U_*\rho_{v,\gamma}$ of $\rho_{v,\gamma}$ under U is $\rho_{Uv,\gamma}$ (since U sends $\Theta(v, \gamma)$ to $\Theta(Uv, \gamma)$ and preserves the uniform distribution π on S^{d-1}). It follows that

$$U A_{v,\gamma} U^{-1} = \int_{\Theta} (U\theta)(U\theta)^\top \rho_{v,\gamma}(\text{d}\theta) = \int_{\Theta} \theta \theta^\top \rho_{Uv,\gamma}(\text{d}\theta) = A_{Uv,\gamma}. \quad (87)$$

In particular, $A_{v,\gamma}$ commutes with every isometry $U \in O(d)$ such that $Uv = v$. Taking U to be the orthogonal reflection with respect to $H_v := (\mathbf{R}v)^\perp$, $A_{v,\gamma}$ preserves $\ker(U - I_d) = \mathbf{R}v$ and is therefore of the form $\phi_1(v, \gamma)vv^\top + C_{v,\gamma}$ where $\phi_1(v, \gamma) \in \mathbf{R}$ and $C_{v,\gamma}$ is a symmetric operator with $C_{v,\gamma}H_v \subset H_v$ and $C_{v,\gamma}v = v$. Next, taking $U = vv^\top + U_v$ where U_v is an arbitrary isometry of H_v , it follows that $C_{v,\gamma}$ commutes on H_v with all isometries U_v , and is therefore of the form $\phi_2(v, \gamma)P_v$, where $P_v = I_d - vv^\top$ is the orthogonal projection on H_v and $\phi_2(v, \gamma) \in \mathbf{R}$. To summarize, we have:

$$A_{v,\gamma} = \phi_1(v, \gamma)vv^\top + \phi_2(v, \gamma)(I_d - vv^\top).$$

Now, the identity (87) shows that, for every $U \in O(d)$ and v, γ , $\phi_1(Uv, \gamma) = \phi_1(v, \gamma)$ and $\phi_2(Uv, \gamma) = \phi_2(v, \gamma)$; hence, these constants do not depend on v and are simply denoted $\phi_1(\gamma), \phi_2(\gamma)$. Defining $\phi(\gamma) := d \cdot \phi_2(\gamma)$ and $\tilde{\phi}(\gamma) := \phi_1(\gamma) - \phi_2(\gamma)$, we therefore have:

$$A_{v,\gamma} = \tilde{\phi}(\gamma)vv^\top + \phi(\gamma) \cdot \frac{1}{d}I_d. \quad (88)$$

Next, observe that

$$\int_{S^{d-1}} \rho_{v,\gamma} \pi(\text{d}v) = \pi; \quad (89)$$

this follows from the fact that the measure π' on the left-hand side of (89) is a probability distribution on S^{d-1} invariant under any $U \in O(d)$, since

$$U_*\pi' = \int_{S^{d-1}} U_*\rho_{v,\gamma}\pi(\text{d}v) = \int_{S^{d-1}} \rho_{Uv,\gamma}\pi(\text{d}v) = \int_{S^{d-1}} \rho_{v,\gamma}(U_*\pi)(\text{d}v) = \int_{S^{d-1}} \rho_{v,\gamma}\pi(\text{d}v) = \pi'.$$

Equation (89), together with Fubini's theorem, implies that

$$\int_{S^{d-1}} A_{v,\gamma} \pi(\text{d}v) = \int_{S^{d-1}} \int_{S^{d-1}} \theta \theta^\top \rho_{v,\gamma}(\text{d}\theta) \pi(\text{d}v) = \int_{S^{d-1}} \theta \theta^\top \pi(\text{d}\theta) =: A.$$

Since A commutes with isometries (by invariance of π), it is of the form cI_d with $c = \text{Tr}(A)/d = (1/d) \int_{S^{d-1}} \|\theta\|^2 \pi(\text{d}\theta) = 1/d$. Plugging (88) into the previous equality, we obtain

$$\frac{1}{d}I_d = \int_{S^{d-1}} \left[\tilde{\phi}(\gamma)vv^\top + \phi(\gamma) \cdot \frac{1}{d}I_d \right] \pi(\text{d}v) = \frac{1}{d}\tilde{\phi}(\gamma)I_d + \frac{1}{d}\phi(\gamma)I_d,$$

so that $\tilde{\phi}(\gamma) = 1 - \phi(\gamma)$. The decomposition (88) then writes:

$$A_{v,\gamma} = (1 - \phi(\gamma))vv^\top + \phi(\gamma) \cdot \frac{1}{d}I_d.$$

Recalling that $F_{v,\gamma}(\Sigma) = \text{Tr}(\Sigma A_{v,\gamma})$, we obtain the desired expression (85) for $F_{v,\gamma}$.

Finally, note that on the one hand,

$$\langle A_{v,\gamma} v, v \rangle = (1 - \phi(\gamma)) \|v\|^2 + \phi(\gamma) \cdot \frac{1}{d} \|v\|^2 = 1 - \frac{d-1}{d} \phi(\gamma),$$

while on the other hand:

$$\langle A_{v,\gamma} v, v \rangle = \int_{S^{d-1}} \langle \theta, v \rangle^2 \rho_{v,\gamma}(d\theta),$$

so that

$$\phi(\gamma) = \frac{d}{d-1} \int_{S^{d-1}} (1 - \langle \theta, v \rangle^2) \rho_{v,\gamma}(d\theta) \geq 0,$$

where we used that $\langle \theta, v \rangle^2 \leq 1$ by the Cauchy-Schwarz inequality.

Now, let α denote the angle between θ and v . We have $\langle \theta, v \rangle = \cos \alpha$ and $\|\theta - v\|^2 = (1 - \cos \alpha)^2 + \sin^2 \alpha = 2(1 - \cos \alpha)$, so that $\langle \theta, v \rangle = 1 - \frac{1}{2} \|\theta - v\|^2$. Since $\rho_{v,\gamma}(d\theta)$ -almost surely, $\|\theta - v\| \leq \gamma$, this implies

$$1 - \langle \theta, v \rangle^2 = 1 - \left(1 - \frac{1}{2} \|\theta - v\|^2\right)^2 = \|\theta - v\|^2 - \frac{1}{4} \|\theta - v\|^4 \leq \gamma^2.$$

Integrating this inequality over $\rho_{v,\gamma}$ yields $\phi(\gamma) \leq d/(d-1)\gamma^2$; this establishes (86).

5.4 Control of the entropy term

We now turn to the control of the entropy term in (83). Specifically, we will show that, for every $v \in S^{d-1}$ and $\gamma > 0$,

$$\text{KL}(\rho_{v,\gamma}, \pi) \leq d \log \left(1 + \frac{2}{\gamma}\right). \quad (90)$$

First, since $d\rho_{v,\gamma}/d\pi = \pi[\Theta(v, \gamma)]^{-1}$ $\rho_{v,\gamma}$ -almost surely, $\text{KL}(\rho_{v,\gamma}, \pi) = \log \pi[\Theta(v, \gamma)]^{-1}$. Now, let $N = N_c(\gamma, S^{d-1})$ denote the γ -covering number of S^{d-1} , namely the smallest $N \geq 1$ such that there exists $\theta_1, \dots, \theta_N \in S^{d-1}$ with

$$S^{d-1} = \bigcup_{i=1}^N \Theta(\theta_i, \gamma). \quad (91)$$

Applying a union bound to (91) and using the fact that $\pi[\Theta(\theta_i, \gamma)] = \pi[\Theta(v, \gamma)]$ yields $1 \leq N\pi[\Theta(v, \gamma)]$, namely

$$\text{KL}(\rho_{v,\gamma}, \pi) \leq \log N. \quad (92)$$

Now, let $N_p(\gamma, S^{d-1})$ denote the γ -packing number of S^{d-1} , which is the largest number of points in S^{d-1} with pairwise distances at least γ . We have, denoting $B^d = \{x \in \mathbf{R}^d : \|x\| \leq 1\}$,

$$N \leq N_p(\gamma, S^{d-1}) \leq N_p(\gamma, B^d) \leq \left(1 + \frac{2}{\gamma}\right)^d, \quad (93)$$

where the first inequality follows from a comparison of covering and packing numbers [Ver18, Lemma 4.2.8], the second one from the inclusion $S^{d-1} \subset B^d$ and the last one from a volumetric argument [Ver18, Lemma 4.2.13]. Combining (92) and (93) establishes (90).

5.5 Conclusion of the proof

First note that, since $\|X'_i\|^2 = \|X_i\|^2 \wedge d \leq d$ for $1 \leq i \leq n$,

$$\text{Tr}(\widehat{\Sigma}'_n) = \frac{1}{n} \sum_{i=1}^n \|X'_i\|^2 \leq d. \quad (94)$$

Putting together the previous bounds (83), (85), (90) and (94), we get with probability $1 - e^{-nu}$, for every $v \in S^{d-1}$, $\gamma \in (0, 1/2]$,

$$\begin{aligned} \alpha' \log \left(\frac{\lambda}{C^2} \right) - \log 2 - \frac{d}{n} \log \left(1 + \frac{2}{\gamma} \right) - u &\leq \lambda F_{v,\gamma}(\widehat{\Sigma}'_n) \\ &= \lambda \left((1 - \phi(\gamma)) \langle \widehat{\Sigma}'_n v, v \rangle + \phi(\gamma) \cdot \frac{1}{d} \text{Tr}(\widehat{\Sigma}'_n) \right) \\ &\leq \lambda \left[(1 - \phi(\gamma)) \langle \widehat{\Sigma}'_n v, v \rangle + \phi(\gamma) \right] \end{aligned}$$

In particular, rearranging, and using the fact that $\phi(\gamma) \leq 1/2$ for $\gamma \leq 1/2$, as well as $\phi(\gamma) \leq \gamma^2$ and $\lambda_{\min}(\widehat{\Sigma}'_n) = \inf_v \langle \widehat{\Sigma}'_n v, v \rangle$, we get with probability $1 - e^{-nu}$,

$$\lambda_{\min}(\widehat{\Sigma}'_n) \geq \frac{2}{\lambda} \left[\alpha' \log \left(\frac{\lambda}{C^2} \right) - \log 2 - \frac{d}{n} \log \left(1 + \frac{2}{\gamma} \right) - u \right] - 2\gamma^2 \quad (95)$$

We first approximately maximize the above lower bound in γ , given λ . Since $\gamma \leq 1/2$, $1 + 2/\gamma \leq 1 + 1/\gamma^2 \leq 5/(4\gamma^2)$. We are therefore led to minimize

$$\frac{2d}{\lambda n} \log \left(\frac{5}{4\gamma^2} \right) + 2\gamma^2$$

over $\gamma^2 \leq 1/4$. Now, let $\gamma^2 = d/(2\lambda n)$, which belongs to the prescribed range if

$$\lambda \geq \frac{2d}{n}. \quad (96)$$

For this choice of γ , the lower bound (95) becomes

$$\begin{aligned} \lambda_{\min}(\widehat{\Sigma}'_n) &\geq \frac{2}{\lambda} \left[\alpha' \log \left(\frac{\lambda}{C^2} \right) - \log 2 - \frac{d}{n} \log \left(\frac{5\lambda n}{2d} \right) - u \right] - \frac{d}{\lambda n} \\ &= \frac{2}{\lambda} \left[\left(\alpha' - \frac{d}{n} \right) \log \lambda - \alpha' \log C^2 - \left\{ \log 2 + \frac{d}{n} \log \left(\frac{5n}{2d} \right) + \frac{d}{2n} \right\} - u \right] \end{aligned}$$

Now, recall that by assumption, $d/n \leq \alpha/6 \leq 1/6$, so that (by monotonicity of $x \mapsto -x \log x$ on $(0, e^{-1}]$, replacing d/n by $1/6$) the term inside the braces is smaller than $c_0 = 1.3$. In addition, assume that $\lambda \geq C^4$, so that $\log(\lambda/C^4) \geq 0$; in this case, condition (96) is automatically satisfied, since $2d/n \leq 1/3 \leq C^4$. Finally, since $\alpha' = \alpha/(2 + \alpha) \geq \alpha/3$ and $d/n \leq \alpha/6$, $\alpha' \leq 2(\alpha' - d/n)$ and $\alpha' - d/n \geq \alpha/6$, so that

$$\left(\alpha' - \frac{d}{n} \right) \log \lambda - \alpha' \log C^2 \geq \left(\alpha' - \frac{d}{n} \right) \log \left(\frac{\lambda}{C^4} \right) \geq \frac{\alpha}{6} \log \left(\frac{\lambda}{C^4} \right),$$

the previous inequalities implies that, for every $\lambda \geq C^4$ and $u > 0$, with probability at least $1 - e^{-nu}$,

$$\lambda_{\min}(\widehat{\Sigma}'_n) \geq \frac{2}{\lambda} \left[\frac{\alpha}{6} \log \left(\frac{\lambda}{C^4} \right) - c_0 - u \right] = \frac{\alpha}{3C^4} \frac{\log \lambda' - 6\alpha^{-1}(c_0 + u)}{\lambda'}$$

where $\lambda' = \lambda/C^4 \geq 1$. A simple analysis shows that for $c \in \mathbf{R}$, the function $\lambda' \mapsto (\log \lambda' - c)/\lambda'$ admits a maximum on $(0, +\infty)$ of e^{-c-1} , reached at $\lambda' = e^{c+1}$. Here $c = 6\alpha^{-1}(c_0 + u) > 0$, so that $\lambda' > e > 1$. Hence, for every $u > 0$, with probability at least $1 - e^{-nu}$,

$$\lambda_{\min}(\widehat{\Sigma}'_n) \geq \frac{\alpha}{3C^4} \exp\left(-1 - \frac{6(c_0 + u)}{\alpha}\right) \geq C'^{-1} e^{-6u/\alpha} =: t, \quad (97)$$

where we let $C' := 3C^4 e^{1+9/\alpha}$ (using the fact that $6c_0 \leq 8$ and $1/\alpha \leq e^{1/\alpha}$). Inverting the bound (97), we obtain that for every $t < C'^{-1}$,

$$\mathbb{P}(\lambda_{\min}(\widehat{\Sigma}'_n) \leq t) \leq (C't)^{\alpha n/6}.$$

Since $\lambda_{\min}(\widehat{\Sigma}_n) \geq \lambda_{\min}(\widehat{\Sigma}'_n)$, and since the bound trivially holds for $t \geq C'^{-1}$, this concludes the proof.

6 Conclusion

In this work, we performed a detailed decision-theoretic analysis of linear prediction under square loss with random design, by providing matching upper and lower bounds on the minimax risk under weak conditions. In particular, we showed that the minimax risk is determined by the distribution of statistical leverage scores, and is approximately minimized in high dimension by centered Gaussian covariates. In particular, we obtained the first upper bounds on the expected risk of the ordinary least squares estimator in the random design setting with non-Gaussian covariates. Those bounds scale as most as $(1 + o(1))\sigma^2 d/n$ as $d = o(n)$ with noise level σ^2 , under some mild conditions on the distribution of covariates.

The previous upper bounds relied on a study of the lower tail and negative moments of empirical covariance matrices. We showed a general lower bound on this lower tail in dimension $d \geq 2$, and established a matching upper bound under a necessary “small-ball” regularity condition on the design. The proof of this result relied on the use of PAC-Bayesian smoothing of empirical processes, with refined non-Gaussian smoothing distributions.

It is worth noting that our upper bound on the lower tail of $\lambda_{\min}(\widehat{\Sigma}_n)$ (Theorem 4) requires $n \geq 6d$; while we did not attempt to optimize the constant factor 6, the approach used here is not sufficient to obtain meaningful bounds for square (or nearly square) design matrices, whose aspect ratio d/n is close to 1. It would be interesting to see if the bound of Theorem 4 can be extended to this case (for instance with centered, variance 1 independent coordinates with bounded density, as in Section 3.3, or more generally under Assumption 1 with $\alpha = 1$), by using the techniques from [RV08, RV09, TV09b, TV09a].

References

- [AC10] Jean-Yves Audibert and Olivier Catoni. Linear regression through PAC-Bayesian truncation. *arXiv preprint arXiv:1010.0072*, 2010.
- [AC11] Jean-Yves Audibert and Olivier Catoni. Robust linear least squares regression. *The Annals of Statistics*, 39(5):2766–2794, 2011.
- [AGZ10] Greg W. Anderson, Alice Guionnet, and Ofer Zeitouni. *An introduction to random matrices*, volume 118 of *Cambridge Studies in Advanced Mathematics*. Cambridge University Press, 2010.

- [ALPTJ10] Radosław Adamczak, Alexander Litvak, Alain Pajor, and Nicole Tomczak-Jaegermann. Quantitative estimates of the convergence of the empirical covariance matrix in log-concave ensembles. *Journal of the American Mathematical Society*, 23(2):535–561, 2010.
- [And03] Theodore W. Anderson. *An Introduction to Multivariate Statistical Analysis*. Wiley New York, 2003.
- [AW01] Katy S. Azoury and Manfred K. Warmuth. Relative loss bounds for on-line density estimation with the exponential family of distributions. *Machine Learning*, 43(3):211–246, 2001.
- [BF83] Leo Breiman and David Freedman. How many variables should be entered in a regression equation? *Journal of the American Statistical Association*, 78(381):131–136, 1983.
- [Bha09] Rajendra Bhatia. *Positive Definite Matrices*, volume 16 of *Princeton Series in Applied Mathematics*. Princeton University Press, 2009.
- [BKM⁺15] Peter L. Bartlett, Wouter M. Koolen, Alan Malek, Eiji Takimoto, and Manfred K. Warmuth. Minimax fixed-design linear regression. In *Proceedings of the 28th Conference on Learning Theory (COLT)*, pages 226–239, 2015.
- [BLM13] Stéphane Boucheron, Gábor Lugosi, and Pascal Massart. *Concentration Inequalities: A Nonasymptotic Theory of Independence*. Oxford University Press, Oxford, 2013.
- [BS10] Zhidong Bai and Jack W. Silverstein. *Spectral Analysis of Large Dimensional Random Matrices*. Springer Series in Statistics. Springer-Verlag New York, 2 edition, 2010.
- [BTW07] Florentina Bunea, Alexandre B. Tsybakov, and Marten H. Wegkamp. Aggregation for gaussian regression. *The Annals of Statistics*, 35(4):1674–1697, 2007.
- [BV04] Stephen Boyd and Lieven Vandenberghe. *Convex Optimization*. Cambridge University Press, 2004.
- [Cat04] Olivier Catoni. *Statistical Learning Theory and Stochastic Optimization: Ecole d’Eté de Probabilités de Saint-Flour XXXI - 2001*, volume 1851 of *Lecture Notes in Mathematics*. Springer-Verlag Berlin Heidelberg, 2004.
- [Cat07] Olivier Catoni. *PAC-Bayesian Supervised Classification: The Thermodynamics of Statistical Learning*, volume 56 of *IMS Lecture Notes Monograph Series*. Institute of Mathematical Statistics, 2007.
- [CDV07] Andrea Caponnetto and Ernesto De Vito. Optimal rates for the regularized least-squares algorithm. *Foundations of Computational Mathematics*, 7(3):331–368, 2007.
- [CH88] Samprit Chatterjee and Ali S. Hadi. *Sensitivity analysis in linear regression*, volume 327 of *Wiley Series in Probability and Statistics*. John Wiley & Sons, New York, 1988.
- [CS02a] Felipe Cucker and Steve Smale. Best choices for regularization parameters in learning theory: on the bias-variance problem. *Foundations of Computational Mathematics*, 2(4):413–428, 2002.

- [CS02b] Felipe Cucker and Steve Smale. On the mathematical foundations of learning. *Bulletin of the American Mathematical Society*, 39(1):1–49, 2002.
- [Dic16] Lee H. Dicker. Ridge regression and asymptotic minimax estimation over spheres of growing dimension. *Bernoulli*, 22(1):1–37, 2016.
- [DM16] David Donoho and Andrea Montanari. High dimensional robust m-estimation: asymptotic variance via approximate message passing. *Probability Theory and Related Fields*, 166(3):935–969, 2016.
- [DVCR05] Ernesto De Vito, Andrea Caponnetto, and Lorenzo Rosasco. Model selection for regularized least-squares algorithm in learning theory. *Foundations of Computational Mathematics*, 5(1):59–85, 2005.
- [DW18] Edgar Dobriban and Stefan Wager. High-dimensional asymptotics of prediction: Ridge regression and classification. *The Annals of Statistics*, 46(1):247–279, 2018.
- [EK13] Nouredine El Karoui. Asymptotic behavior of unregularized and ridge-regularized high-dimensional robust regression estimators: rigorous results. *arXiv preprint arXiv:1311.2445*, 2013.
- [EK18] Nouredine El Karoui. On the impact of predictor geometry on the performance on high-dimensional ridge-regularized generalized robust regression estimators. *Probability Theory and Related Fields*, 170(1):95–175, 2018.
- [EKK11] Nouredine El Karoui and Holger Kösters. Geometric sensitivity of random matrix results: consequences for shrinkage estimators of covariance and related statistical methods. *arXiv preprint arXiv:1105.1404*, 2011.
- [Ess66] Carl G. Esseen. On the Kolmogorov-Rogozin inequality for the concentration function. *Zeitschrift für Wahrscheinlichkeitstheorie und Verwandte Gebiete*, 5(3):210–216, 1966.
- [Fed96] Herbert Federer. *Geometric measure theory*. Springer, 1996.
- [Fos91] Dean P. Foster. Prediction in the worst case. *The Annals of Statistics*, 19:1084–1090, 1991.
- [GCS⁺13] Andrew Gelman, John B. Carlin, Hal S. Stern, David B. Dunson, Aki Vehtari, and Donald B. Rubin. *Bayesian Data Analysis*. Chapman and Hall/CRC, 2013.
- [GKKW02] László Györfi, Michael Kohler, Adam Krzyżak, and Harro Walk. *A distribution-free theory of nonparametric regression*. Springer Science & Business Media, 2002.
- [HJ90] Roger A. Horn and Charles R. Johnson. *Matrix Analysis*. Cambridge University Press, Cambridge, 1990.
- [HKZ14] Daniel Hsu, Sham M. Kakade, and Tong Zhang. Random design analysis of ridge regression. *Foundations of Computational Mathematics*, 14(3):569–600, 2014.
- [HMRT19] Trevor Hastie, Andrea Montanari, Saharon Rosset, and Ryan J. Tibshirani. Surprises in high-dimensional ridgeless least squares interpolation. *arXiv preprint arXiv:1903.08560*, 2019.

- [Hoe62] Arthur E. Hoerl. Application of ridge analysis to regression problems. *Chemical Engineering Progress*, 58:54–59, 1962.
- [HS16] Daniel Hsu and Sivan Sabato. Loss minimization and parameter estimation with heavy tails. *Journal of Machine Learning Research*, 17(18):1–40, 2016.
- [Hub73] Peter J. Huber. Robust regression: asymptotics, conjectures and Monte Carlo. *The Annals of Statistics*, 1(5):799–821, 1973.
- [Hub81] Peter J. Huber. *Robust statistics*. John Wiley and Sons, 1981.
- [HW78] David C. Hoaglin and Roy E. Welsch. The hat matrix in regression and ANOVA. *The American Statistician*, 32(1):17–22, 1978.
- [KM15] Vladimir Koltchinskii and Shahar Mendelson. Bounding the smallest singular value of a random matrix without concentration. *International Mathematics Research Notices*, 2015(23):12991–13008, 2015.
- [Led01] Michel Ledoux. *The Concentration of Measure Phenomenon*, volume 89. American Mathematical Society, Providence, RI, 2001.
- [LM16] Guillaume Lecué and Shahar Mendelson. Performance of empirical risk minimization in linear aggregation. *Bernoulli*, 22(3):1520–1534, 2016.
- [LM19] Gábor Lugosi and Shahar Mendelson. Mean estimation and regression under heavy-tailed distributions: a survey. *Foundations of Computational Mathematics*, 19:1145–1190, 2019.
- [LST03] John Langford and John Shawe-Taylor. PAC-Bayes & margins. In *Advances in Neural Information Processing Systems 15*, pages 439–446, 2003.
- [McA99a] David A. McAllester. PAC-Bayesian model averaging. In *Proceedings of the 12th Annual Conference on Computational Learning Theory (COLT)*, pages 164–170, 1999.
- [McA99b] David A. McAllester. Some PAC-Bayesian theorems. *Machine Learning*, 37(3):355–363, 1999.
- [Men15] Shahar Mendelson. Learning without concentration. *Journal of the ACM*, 62(3):21, 2015.
- [MP14] Shahar Mendelson and Grigoris Paouris. On the singular values of random matrices. *Journal of the European Mathematical Society*, 16(4):823–834, 2014.
- [Nem00] Arkadi Nemirovski. Topics in non-parametric statistics. *Lectures on Probability Theory and Statistics: Ecole d’Ete de Probabilites de Saint-Flour XXVIII-1998*, 28:85–277, 2000.
- [NV13] Hoi H. Nguyen and Van H. Vu. Small ball probability, inverse theorems, and applications. In *Erdős Centennial*, pages 409–463. Springer, 2013.
- [Oli16] Roberto Imbuzeiro Oliveira. The lower tail of random quadratic forms with applications to ordinary least squares. *Probability Theory and Related Fields*, 166(3):1175–1194, 2016.

- [RM16] Garvesh Raskutti and Michael W. Mahoney. A statistical perspective on randomized sketching for ordinary least-squares. *Journal of Machine Learning Research*, 17(1):7508–7538, 2016.
- [Rog87] Boris A. Rogozin. The estimate of the maximum of the convolution of bounded densities. *Teoriya Veroyatnostei i ee Primeneniya*, 32(1):53–61, 1987.
- [RV08] Mark Rudelson and Roman Vershynin. The Littlewood–Offord problem and invertibility of random matrices. *Advances in Mathematics*, 218(2):600–633, 2008.
- [RV09] Mark Rudelson and Roman Vershynin. Smallest singular value of a random rectangular matrix. *Communications on Pure and Applied Mathematics*, 62(12):1707–1739, 2009.
- [RV10] Mark Rudelson and Roman Vershynin. Non-asymptotic theory of random matrices: extreme singular values. In *Proceedings of the International Congress of Mathematicians*, volume III, pages 1576–1602, 2010.
- [RV14] Mark Rudelson and Roman Vershynin. Small ball probabilities for linear images of high-dimensional distributions. *International Mathematics Research Notices*, 2015(19):9594–9617, 2014.
- [Sha15] Ohad Shamir. The sample complexity of learning linear predictors with the squared loss. *Journal of Machine Learning Research*, 16(108):3475–3486, 2015.
- [SHS09] Ingo Steinwart, Don Hush, and Clint Scovel. Optimal rates for regularized least squares regression. In *Proceedings of the 22nd Annual Conference on Learning Theory (COLT)*, pages 79–93, 2009.
- [Ste60] Charles Stein. Multiple regression. In *Contributions to Probability and Statistics: Essays in Honor of Harold Hotelling*. Stanford University Press, 1960.
- [SV13] Nikhil Srivastava and Roman Vershynin. Covariance estimation for distributions with $2 + \varepsilon$ moments. *The Annals of Probability*, 41(5):3081–3111, 2013.
- [SZ07] Steve Smale and Ding-Xuan Zhou. Learning theory estimates via integral operators and their approximations. *Constructive Approximation*, 26(2):153–172, 2007.
- [Tao12] Terence Tao. *Topics in random matrix theory*, volume 132 of *Graduate Studies in Mathematics*. American Mathematical Society, 2012.
- [Tik63] Andrey N. Tikhonov. Solution of incorrectly formulated problems and the regularization method. *Soviet Mathematics Doklady*, 4:1035–1038, 1963.
- [Tsy03] Alexandre B. Tsybakov. Optimal rates of aggregation. In *Learning Theory and Kernel Machines*, volume 2777 of *Lecture Notes in Artificial Intelligence*, pages 303–313. Springer Berlin Heidelberg, 2003.
- [Tsy09] Alexandre B. Tsybakov. *Introduction to nonparametric estimation*. Springer, 2009.
- [TV04] Antonia M. Tulino and Sergio Verdú. Random matrix theory and wireless communications. *Foundations and Trends® in Communications and Information Theory*, 1(1):1–182, 2004.

- [TV09a] Terence Tao and Van Vu. From the Littlewood-Offord problem to the circular law: universality of the spectral distribution of random matrices. *Bulletin of the American Mathematical Society*, 46(3):377–396, 2009.
- [TV09b] Terence Tao and Van H. Vu. Inverse Littlewood-Offord theorems and the condition number of random discrete matrices. *Annals of Mathematics*, 169(2):595–632, 2009.
- [Ver12] Roman Vershynin. *Introduction to the non-asymptotic analysis of random matrices*, pages 210–268. Cambridge University Press, Cambridge, 2012.
- [Ver18] Roman Vershynin. *High-dimensional probability: An introduction with applications in data science*, volume 47. Cambridge University Press, 2018.
- [Vov01] Volodya Vovk. Competitive on-line statistics. *International Statistical Review*, 69(2):213–248, 2001.
- [Was06] Larry Wasserman. *All of Nonparametric Statistics*. Springer Texts in Statistics. Springer-Verlag New York, Inc., Secaucus, NJ, USA, 2006.
- [Yas14] Pavel Yaskov. Lower bounds on the smallest eigenvalue of a sample covariance matrix. *Electronic Communications in Probability*, 19, 2014.
- [Yas15] Pavel Yaskov. Sharp lower bounds on the least singular value of a random matrix without the fourth moment condition. *Electronic Communications in Probability*, 20, 2015.